
Less is More: Early Stopping Rollout for On-Policy Distillation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 On-policy distillation has recently emerged as a promising alternative to standard
2 sequence-level imitation, training a student by scoring its own rollouts with a
3 teacher model. However, one important limitation is observed in this paradigm:
4 at each position, the teacher scores the student’s next token conditioned on the
5 student’s own previous outputs. Early tokens are therefore judged mainly in the
6 context of the original problem and reflect how the student approaches the problem,
7 whereas later tokens are constrained by increasingly student-specific context and
8 provide much weaker supervision. Based on this intuition, we propose Early
9 Stopping Rollout (ESR), a simple yet effective distillation strategy that restricts
10 both rollout generation and reverse-KL to the first N response tokens. This position-
11 focused objective emphasizes the tokens that encode planning and reasoning, rather
12 than spending supervision budget on late tokens whose informational value rapidly
13 degrades. Empirical experiments have proved that ESR shows both surpassing
14 performance with high efficiency under various models and training methods. It
15 is also proved that the teacher collapses in sufficiently long rollouts suggesting
16 that late-position supervision contributes little useful learning signal. We further
17 investigate the mechanism behind this effect and show that position is a distinct
18 and crucial dimension of supervision, not merely a proxy for token uncertainty or
19 difficulty. It establishes position-aware supervision as an important direction for
20 understanding and advancing OPD in future studies.

21 1 Introduction

22 On-policy distillation (OPD) has emerged as a dominant paradigm for model distillation in industrial
23 practice. The student generates its own rollouts τ , which are then scored by the teacher: at each token,
24 the teacher’s probability $\pi_{teacher}(\tau_{student} \mid x_{prompt})$ serves as the soft target for the student [Agarwal
25 et al., 2024, Gu et al., 2024]. Viewed through an RL lens, this is equivalent to using the teacher as a
26 dense, token-level reward model that judges the student’s own behavior on a given prompt. The final
27 objective is then a simple *uniform average* of the per-token reverse-KL across the trajectory.

28 This uniform averaging implicitly assumes that the teacher’s signal is equally reliable at every position
29 in the trajectory. We argue that this assumption is unfounded, and that it breaks down precisely
30 because of the autoregressive structure of LLMs. At the first token, the teacher’s score is conditioned
31 only on the prompt: $\pi_{teacher}(\tau_{student}^{t=1} \mid x_{prompt})$ - indeed what we expect the teacher to score on. At
32 a later position m , however, it becomes conditioned also on the student’s own previously generated
33 tokens: $\pi_{teacher}(\tau_{student}^{t=m} \mid x_{prompt}, \tau_{student}^{t=1:m})$. This conditioning context is off-policy to the teacher
34 model and may even be out-of-distribution (OOD), and grows more so with position.

35 To probe how severe this drift is, we run a preliminary experiment that measures how the teacher
36 behaves when continuing from an early-stopped student rollout. As shown in Figure 1, the teacher’s

quality collapses toward the student’s within roughly 300 tokens. If the teacher’s late-position signal is unreliable, the natural question is whether early tokens alone are systematically more informative — and whether training on them is sufficient.

Motivated by this, we propose Early Stopping Rollout(ESR): restrict the student rollout to its first N tokens (e.g., 50 or 100) and compute the distillation loss only on this early window. The change is a single line in any on-policy distillation loop. Despite its simplicity, ESRconsistently outperforms full-rollout OPD across tasks (math, code, function calling), training regimes (LoRA, full fine-tuning), model scales (students 1.5B–7B, teachers 1.7B–14B), and model families (Qwen, Gemma), while reducing wall-clock cost by up to $24\times$ and peak training memory by up to $4\times$. We further observe that full-rollout OPD frequently destabilizes and collapses during training, whereas ESRdoes not collapse in any of our settings. Perhaps most surprisingly, although the student is normally upper-bounded by the teacher’s performance, ESR-trained students often *exceed* the teacher.

Such a small change yields such large gains probe us to investigate deeper. Firstly, we want to further validate whether the effectiveness of earlier tokens is intrinsically caused by better KL and entropy signals, because we also find that early tokens generally have larger KL and entropy. We test selectors that pick tokens by high KL, high entropy or combined regardless of position; all underperform ESR, with most underperforming full-rollout. Therefore, although we don’t exclude KL and entropy as auxillary factors, we confirm that the position is a an independent factor rather a saliency proxy. Furthermore, we identified an additional important mechanism *Convergence Cascade*: after training on the early window, KL divergence on the *untrained* late tokens also drops by 30–40%. Therefore, the KL divergence does not have to see late positions to repair them.

Lastly, we investigated into why ESRsometimes even *exceeds* the teacher — a stronger phenomenon we observe on most of our main experiment Table 1. We trace this to the mode-seeking nature of reverse-KL: the student commits to a subset of supported teacher modes instead of chasing all of them. Our ESRloss helps it steer toward a non-dominant but correct mode rather than the teacher’s argmax. Empirically, we find ESR-trained students commit harder to their own choices than full-rollout students yet agree *less* with the teacher’s argmax, and produce trajectories $\sim 3\times$ shorter than full-rollout and $\sim 2\times$ shorter than the base student (§5.3). It learns the teacher’s strategy but preserves its succinct style, leading to better compound behaviors. This finding showcases a potential path of superceding the teacher model in distillation that worth further investigations.

Our contributions are summarized below:

1. **Early Stopping Rollout(ESR) outperforms full-rollout on-policy distillation.** A one-line change—restricting the rollout length to the first N response tokens—beats full-sequence distillation across tasks, model families, scales, and training regimes, while being dramatically more efficient and far more stable to train. It can even exceed the teacher’s own performance (Sections 3–4.3).
2. **A systematic investigation of why it works.** First, through a token-selection ablation, we show that the gain is not “high KL” or “high entropy” in disguise: position itself is a load-bearing axis. We also identify convergence cascading to be another important mechanism. Furthermore, we explain why ESRoften surpasses the teacher: by exploiting the mode-seeking behavior of reverse-KL training, ESRpushes the student toward a better, non-dominant sub-mode (Section 5).

2 Why Late-Position Supervision Is Unreliable

In on-policy distillation, the teacher model scores the student’s rollout at each position t via $\pi_t(\cdot \mid x, y_{<t}^{\text{student}})$ and averages these scores uniformly across the rollout. The conditioning context $y_{<t}^{\text{student}}$, however, is whatever this particular student happened to write at this particular moment of training. As t grows it drifts away from any distribution the teacher itself was trained on—off-policy to the teacher and, for long enough or atypical enough rollouts, plausibly out-of-distribution. We term this problem as “Off-policy Teacher Drift”.

To verify the intuition, we feed the teacher a k -token student-generated rollout on MATH-500 and letting it continue autoregressively, the teacher’s avg@4 ($n=4$, $t=0.7$) falls from its 65.30% unconditional baseline to 62.70% at $k=100$ and 51.75% at $k=300$ —essentially student-baseline

(Figure 1). Late-position teacher scores are not independent assessments of the prompt; it is a guess at how to continue a trajectory the student has already committed to.

The problem compounds during training because the student’s policy is moving. Early on, the student’s rollout is far from the teacher; late-position targets are fitted under a strongly off-policy prefix. As training progresses and earlier tokens approach the teacher’s distribution, the right late-position behavior changes too—continuations that were sensible above an off-policy prefix are no longer sensible above an on-policy one. But the late-position habits learned earlier do not cleanly erase: gradient descent gives no guarantee of unlearning, and in LoRA’s low-rank setting it updates make redirection compete for capacity that may already be spent. Full-rollout training thus chases a moving late-position target while carrying ballast from earlier ones. We term this problem as “Stale Late-position Lock-in”. This insight expects a training stability issues for full-rollout OPD – a prediction we later verify in §4.3

3 Our Method: Early Stopping Rollout (ESR) On Policy Distillation

3.1 Formulation

Let π_s denote the student and π_t the teacher. In standard on-policy reverse-KL distillation, the student generates a response $\mathbf{y} = (y_1, \dots, y_T)$ conditioned on prompt x , and the loss is

$$\mathcal{L}_{\text{full}} = \mathbb{E}_{\mathbf{y} \sim \pi_s(\cdot|x)} \left[\sum_{t=1}^T \text{KL}(\pi_s(\cdot | x, \mathbf{y}_{<t}) \parallel \pi_t(\cdot | x, \mathbf{y}_{<t})) \right]. \quad (1)$$

ESR (position cutoff N , with $N \ll T$ in practice) truncates the student rollout to its first N tokens, and the loss is computed over exactly those tokens:

$$\mathcal{L}_{\text{ESR}}(N) = \mathbb{E}_{\mathbf{y} \sim \pi_s(\cdot|x), |\mathbf{y}| \leq N} \left[\sum_{t=1}^{\min(N, |\mathbf{y}|)} \text{KL}(\pi_s(\cdot | x, \mathbf{y}_{<t}) \parallel \pi_t(\cdot | x, \mathbf{y}_{<t})) \right]. \quad (2)$$

If the student emits EOS before position N , the rollout terminates naturally. Everything else—generation temperature, LoRA target modules, optimizer, scorer—is unchanged from the standard on-policy KD loop.

4 Analysis on the Advancement of ESR

4.1 Setup

Models. We evaluate across multiple student-teacher pairs. Our primary student is Qwen2.5-Math-1.5B [Qwen Team, 2024], paired with teachers of increasing scale: Qwen3-1.7B, Qwen3-4B, and Qwen3-8B [Qwen Team, 2025]. We also tested Qwen2.5-Math-7B with Qwen3-14B to test student model size generality. To test cross-family generality, we also evaluate with gemma-2-2b-it (student) and gemma-3-4b-it (teacher).

Training. All experiments use LoRA [Hu et al., 2022] ($r=32$, $\alpha=64$, all linear layers, $\text{lr} = 5 \times 10^{-5}$) with reverse KL divergence loss. Temperature 0.7 for on-policy generation. Each training step processes a batch of 16 problems with 1 rollout per problem ($n_{\text{samples}}=1$, batch size 16). We train for 200 steps on all tasks, and saving checkpoints every 50 steps. Training data are drawn from NuminaMath, CodeUltraFeedback, and glaive-function-calling-v2. Our method uses $N=100$ unless otherwise specified.

Evaluation. MATH-500 [Hendrycks et al., 2021, Lightman et al., 2023] with $n=4$ samples at temperature 0.7 (reporting avg@4), HumanEval/HumanEval+ [Chen et al., 2021, Austin et al., 2021, Liu et al., 2023] at temperature 0.0 (reporting pass@1), BFCL [Patil et al., 2025] with 600 problems reporting full accuracy: correct function name and arguments.

4.2 Overall performance

Table 1 presents the main results across tasks and model families. Table 2 presents ablation results with full finetuning. Table 2 shows model size scaling results. We report the best performance

Table 1: **Main results: ESR dominates Full Rollout across tasks and families.** Best score across training steps; metrics defined in §4 setup. ★: ESR surpasses the teacher reference. **Full Rollout** numbers mark the best step before degradation; ↓ flags cells where Full Rollout further decays with training (more details in Appendix C).

Pair	Method	MATH-500		Coding		BFCL
		avg@4	pass@4	HE	HE+	full_acc
Qwen2.5-Math 1.5B → Qwen3 1.7B	Student baseline	50.95	72.80	31.10	26.20	2.70
	Teacher reference	65.30	77.00	39.60	36.00	54.00
	Full Rollout	62.35↓	74.60↓	40.20↓	35.40↓	58.20★
	ESR	65.85★	79.80★	42.10★	36.00	61.30★
Gemma-2 2B → Gemma-3 4B	Student baseline	13.45	28.20	23.78	20.10	73.17
	Teacher reference	66.60	74.80	20.70	20.10	72.83
	Full Rollout	22.95	31.40	22.00↓	17.70↓	76.83★
	ESR	27.20	39.40	28.70★	24.40★	79.00★

Table 2: **FullFT and LoRA scaling on MATH-500.** (Left) FullFT for Qwen and Gemma student/teacher pairs. (Right) LoRA scaling across teacher and student sizes; ↓ marks the **Full Rollout** peak before degradation. Both: best across training steps; **bold**: ESR beats **Full Rollout**.

Pair	Method	avg@4	pass@4
Qwen2.5-Math 1.5B	student / teacher	50.95 / 65.30	72.80 / 77.00
→ Qwen3 1.7B	Full Rollout	58.20	75.40
	ESR	56.20	73.80
Gemma-2 2B → Gemma-3 4B	student / teacher	13.45 / 66.60	28.20 / 74.80
	Full Rollout	13.90	25.00
	ESR	26.65	40.40

Pair	Method	avg@4	pass@4
Qwen2.5-Math 1.5B	student / teacher	50.95 / 77.95	72.80 / 86.40
→ Qwen3 4B	Full Rollout	67.45↓	80.60
	ESR	68.95	81.00
Qwen2.5-Math 1.5B	student / teacher	50.95 / 77.70	72.80 / 87.80
→ Qwen3 8B	Full Rollout	66.75	81.60
	ESR	67.85	82.20
Qwen2.5-Math 7B	student / teacher	53.60 / 76.15	75.00 / 83.20
→ Qwen3 14B	Full Rollout	68.85↓	80.00
	ESR	68.95	81.20

130 achieved across all training steps for each configuration. But notice that full rollout distillation
131 frequently collapse to much lower performance as it trains. We denote the sessions that collapses
132 with a down arrows.

133 **ESRbeats full rollout across model families and tasks, and frequently surpasses teacher per-**
134 **formance.** Across the cells of Table 1, ESR beats full rollout on every cell, and in six cells ESR
135 *surpasses the teacher* outright (Qwen math avg@4, pass@4; Qwen HumanEval and HumanEval+;
136 Qwen BFCL full_acc; Gemma coding on all four columns; Gemma BFCL full_acc). The Gemma
137 coding results have the biggest gap: full rollout drives HumanEval pass@1 to 22.00, *below* the 23.78
138 student baseline and worse than the 20.70 teacher, while ESR recovers 31.10. Qwen-math full rollout
139 is marked ↓ because its best number (62.35 avg@4) is a non-stable best step: under $n=16$ it collapses
140 to 46.3% by step 100 (Appendix C.6).

141 **ESRmatches or beats full rollout across full fine-tuning and LoRA settings.** Table 2 reports
142 FullFT on MATH-500 for both Qwen and Gemma. On Qwen **Full Rollout** 58.20 is slightly better
143 than ours ESR 56.20 avg@4, but the gap is close. For Gemma ESR still dominates **Full Rollout**
144 by +12.75pp avg@4 and +15.40pp pass@4. ESR is therefore the safer choice in both parameter
145 regimes.

146 **ESRscales across model sizes.** The advantage persists as both teacher and student grow. Table 2
147 reports the results scaling both the teacher and the student. With the same student (Qwen2.5-Math-
148 1.5B), we tried Qwen3-4B and 8B as the teacher. With Qwen3-4B, **Full Rollout** peaks at (67.45) and
149 then *collapses* to 55.05 (Appendix C.6), while ESR stays stable and climbs to 68.95. Against the 8B
150 teacher, ESR still edges **Full Rollout** (67.85 vs 66.75). The pattern persists at larger scale along with
151 student model too: with a Qwen2.5-Math-7B student and Qwen3-14B teacher, **Full Rollout** again
152 peaks at step 50 (68.85) and drifts to 62.40, while ESR holds stable across all checkpoints.

Figure 1: **(Left) Off-policy teacher drift.** The teacher loses accuracy quickly as the student-generated rollout grows over a few hundred tokens. MATH-500, avg@4 ($n=4$, $t=0.7$). Teacher = Qwen3-1.7B; student = Qwen2.5-Math-1.5B. After ~ 300 student tokens the teacher has effectively been dragged down to student-baseline performance. **(Right) Rollout length N sweep on MATH-500.** LoRA, Qwen2.5-Math-1.5B $\rightarrow$ Qwen3-1.7B; best avg@4 across training steps. Full rollout and the undistilled baseline are shown as horizontal references. Performance saturates for $N \in [50, 200]$ and all beat full-seq.




Table 3: **Training efficiency.** Single A6000 (48 GB), bs=16, student teacher Qwen3-1.7B. We report the average running time and GPU memory usage across student model generation, training and teacher scoring phases. Note that the real time usage can be larger if there isn’t enough GPU to hold the student and teacher models together and requires model loading and unloading.

Phase	Per-step wall time			Peak GPU memory		
	ESR (100 tok.)	Full rollout	Speedup	ESR (100 tok.)	Full rollout	Savings
Teacher scoring	1 s	7 s	7×	7.3 GB	14.9 GB	2.0×
Student generation	5 s	180 s	36×	7.2 GB	8.9 GB	1.2×
Student training	2 s	7 s	3.5×	9.6 GB	39.5 GB	4.1×
Total (single step)	8 s	194 s	24×	24.1 GB	63.3 GB	2.6×
Total (200 steps)	27 min	10.8 hr	24×			

4.3 Stability of ESR

ESR is more robust than full rollout in training. Across our (pair $\times$ task $\times$ regime) matrix, full-rollout distillation degrades with training very frequently; ESR degrades nowhere. Appendix Table 12 reports each collapsing cell as a (peak $\rightarrow$ lowest) pair with the corresponding step indices. For example, Qwen-LoRA coding falls from 40.2 to 26.8; Qwen-LoRA math collapses from 65.60 at s50 to 43.80 ($\Delta = -21.80$ pp, ending below the 50.95% student baseline); Gemma-LoRA collapse from 22.0 to 11.6 ($\Delta = -10.4$ pp). Three of seven full-rollout collapses end below the student baseline. ESR generally hold a fluctuation within ± 3 pp of its peak across training (Tables 7, 9).

Early Stopping Rollouts not sensitive to the choice of N . A natural question naturally occurs – how to choose where to stop? Is it sensitive? We conducted a set of sweeping experiments in Figure 1 sweeps N on MATH-500 and reveals a robust region: it reaches just as good performance starting from $N=50$ and remains stable to $N=200$. The method is not sensitive to the exact choice of N within a certain region.

4.4 Efficiency of ESR

Table 3 shows that ESR achieves a $24\times$ wall-clock speedup and reduces peak training memory by $\sim 4\times$. The dominant cost in full-seq is autoregressive generation (180 s/step for sequences averaging ~ 1000 tokens); ESR generates only $N=100$ tokens (5 s/step). Note that with ESR, all the student and teacher models can be put in one A6000 GPU comfortably. In our own practice, it saves a further big time overhead of model loading and unloading that we do not report here.

Figure 2: **(Left) Case study: early tokens encode strategy; later tokens mostly execute.** A representative MATH-500 trajectory. The first 100 tokens establish the geometric setup, name the unknown, and identify the key relationship (altitude bisects the leg)—these are the choices that determine whether the rollout will succeed. The last 100 tokens carry out an algebraic computation that any solver can finish once the strategy is fixed. ESR supervises the first window; the second follows for free. **(Right) Token selection strategies on MATH-500.** Qwen2.5-Math-1.5B $\rightarrow$ Qwen3-1.7B, LoRA, $n_{\text{samples}}=1$, best avg@4 across training steps. All selectors operate on full-length rollouts; ESR is the only one that additionally truncates the rollout. Baseline (no distillation) 50.95%; full-seq 62.35%. H_s : student entropy; H_t : teacher entropy; RKL: reverse KL.

Problem. In an isosceles right triangle, the altitude to the hypotenuse has length $4\sqrt{2}$. What is the area?

First 100 tokens (strategy). “Let’s start by understanding the problem. In an isosceles right triangle, the two legs are equal, and the hypotenuse is $\sqrt{2}$ times the leg. If we denote each leg by a , then hypotenuse $= a\sqrt{2}$. The altitude to the hypotenuse is half the leg, because the altitude bisects the...”

Last 100 tokens (execution). “...altitude from the right-angle vertex to the hypotenuse. This altitude has length $h = \frac{a \cdot a}{a\sqrt{2}} = \frac{a^2}{a\sqrt{2}} = \frac{a}{\sqrt{2}} = \frac{a\sqrt{2}}{2}$. Setting this equal to $4\sqrt{2}$: $\frac{a\sqrt{2}}{2} = 4\sqrt{2}$, so $a\sqrt{2} = 8\sqrt{2}$, giving $a = 8$. The area is $\frac{1}{2} \times 8 \times 8 = \boxed{32}$.”

Selection method	k	Best avg@4
ESR-100 (ours)	100	65.85
Full-seq (all tokens)	—	62.35
Top-RKL	100	53.35
Top- H_t (teacher entropy)	100	63.30
Top- H_s (student entropy)	100	62.70
Top- H_s	200	53.20
RKL $\cdot H_s$	100	56.90
$H_t \cdot H_s$ (product)	100	55.35
RKL $\cdot H_t \cdot H_s$ (triple product)	100	57.90

5 Additional Reasons Why Does ESR Work?

The surprising effectiveness of Early Stopping Rollout probes us to investigate deeper beyond what we have identified Section 2. We structure this section around three questions: (§5.1) Is the early-position effect induced by high-KL or high-entropy tokens? (§5.2) If not, what additional structural reasons make the early window privileged? (§5.3) Why does ESR sometimes *exceed* the teacher?

5.1 What Does The Effectiveness of ESR Stem From?

In Section 2, we identified the “Off-policy Teacher Drift” and “Stale Late-position Lock-in” problem that can cause the late-position tokens to provide degrading signals. During the experiment analysis, we find that, as shown in Figure 3, early positions simultaneously have high KL divergence between student and teacher model, and high token entropy from both student and teacher models. This finding probes us to wonder if the effectiveness is intrinsically induced by the KL and entropy. To control these factors, we conduct a series of ablation experiments: pick the same amount (100) of tokens based on the highest KL divergence, highest student/teacher entropy or with them in combination with ESR, regardless of position (Figure 2). If the effectiveness is indeed induced by the KL or entropy, then they should reproduce the same or better results.

To our surprise, all underperform ESR, and most of them also much underperform the full rollout results. Teacher or student entropy based selection can match the full sequence by them alone, but their combination falls short significantly. What’s also interesting is that KL divergence measure, the direct calculation of the loss magnitude, barely works. It only improves the baseline (50.95%) for about 3 percent. And more surprisingly, we find that the largest 100 tokens of KL occupies around 93% of the entire trajectory loss. This shows that the tokens that has larger signals are not necessarily the ones that have effective signals.

Therefore, although we still don’t exclude KL and entropy as potential mediator factor, we exclude them to be the sole factors that causes early tokens to be special. Position, therefore, should be considered an independent token selection dimension for the future.

5.2 The Convergence Cascade Effect of ESR

We find an additional important mechanism that assists in ESR’s effectiveness. Even we only trained the first N tokens with ESR, training, per-position KL divergence beyond $[0, N]$ region still drops by 30–40% (Figure 3). This shows that the student can pick up the teacher’s “global mindset” even with just the beginning tokens. One reason that we suspect is that the beginning tokens often consist of problem framing and strategic planning content. The case study in Figure 2 (Left) illustrates

Figure 3: **(Left)** Early tokens are high on student entropy, teacher entropy, and KL divergence simultaneously. **(Right)** **The convergence cascade.** Per-position KL between distilled student and teacher, before vs. after ESRtraining. Yellow band: positions $[0, 100]$ that actually receive training loss. Blue band: the KL gap closed by training. Positions $100+$ —which see no direct training signal—drop to the same KL as the trained region, confirming that alignment on the early window cascades through the autoregressive rollout.





Figure 4: **Mode-seeking on the planning region.** Schematic of reverse-KL behaviour at a multi-modal planning position. The teacher supports two modes (e.g. a verbose plan and a concise correct plan); reverse KL penalises student mass outside the support but not concentration within it. A short early-window loss can collapse the student onto the better supported mode, allowing the student to exceed the teacher’s average behaviour. Full-sequence training reverts the student toward the averaged teacher across late positions, undoing the concentration.



203 this concretely: on a representative MATH-500 trajectory, the first 100 tokens set up the geometry,
 204 name the unknown, and identify the key relationship (the altitude bisects the leg)—the choices that
 205 determine whether the rollout will succeed—while the last 100 tokens merely execute the algebra
 206 that any solver can finish once the strategy is fixed. Therefore once the student picks up how to frame
 207 problems and plan the strategy, the later content naturally follows.

208 Moreover, recently Cloud et al. [2025] shows student models may be able to learn the teacher’s deep
 209 internal preference even with random numbers generated by the teacher, called “subliminal learning”.
 210 Therefore, the early tokens may suffice to inject the subliminal mindset to the student.

211 5.3 Why Does ESR Even Exceed the Teacher?

212 ESR exceeds the teacher in most of the main experiments¹. Even the full rollout model slightly
 213 exceeds the teacher in function calling experiments, although weaker than our method. This shows
 214 that student has the potential exceed the teacher even in normal OPD, and our method amplifies it.
 215 Why is so? Isn’t teacher supposed to be the upper bound?

216 We propose the reason lies in the mechanism of reverse KL $KL(\pi_s || \pi_t)$, which has a mode-seeking
 217 behavior: it penalizes the student for putting mass on tokens the teacher does not support, but not for
 218 concentrating mass on a single supported token. Therefore, the student has the possibility to land on
 219 a submode of the teacher that is actually better. We visualize this mechanism in Figure 4, and verify
 220 it empirically below.

221 **Behavioral signature.** To verify this, we first scan the behavioral differences across distilled models.
 222 Surprisingly, ESR-trained students produce sequences $2-3\times$ shorter than the teacher, full-rollout,

Table 4: **Behavioral Modes Analysis.** *Left:* Response length distribution on MATH-500. ESR learns an output style that dramatically shorter than full rollout, teacher, and even its own baseline model. We show the length of the 10 percentile, 50 percentile (media), 90 percentile and the average length. *Right:* How student’s output aligns with the teacher’s distribution modes. “Top-1 Prob” shows the probability of the top 1 token’s probability mass, indicating how much the student commits in its first choice. “=tch. top-1” means the percentage of times where the student’s first choice is same as that of the teacher. “ $\in$ tch. top 2–5” means the percentage of times where the student’s first choice is the teachers 2nd to 5th choice. And “ $\notin$ tch. top-5” are the times the student’s first choice isn’t in the teacher’s top 5 list. Our method tends to choose teacher’s non-first choice and is more confident in its choice than the full rollout version.

Model	Response length (tokens)					Student commitment & mode alignment			
	10th ptile	50th ptile	90th ptile	Mean	Ratio	Top-1 prob	= tch. top-1	$\in$ tch. top 2–5	$\notin$ tch. top-5
Base student	~400	~860	~1,500	~990	2.2×	0.71	28.6%	59.8%	11.5%
Teacher (Q3-1.7B)	~700	~1,150	~1,800	~1,210	2.6×	—	—	—	—
Full Rollout	~1,190	~1,530	~1,770	~1,480	3.2×	0.77	45.7%	44.6%	9.7%
ESR	~190	~380	~800	~460	1.0×	0.79	41.9%	47.4%	10.7%

and even the base student itself: ESR-100’s median length is ~ 380 tokens, against $\sim 1,150$ for the teacher and $\sim 1,530$ for full-rollout (Table 4, left). The teacher is substantially more verbose than the student, so distilling from such a teacher generally drags the student’s length up — and indeed full-rollout produces rollouts even longer than the teacher itself. The fact that ESR learns from the *same* teacher yet moves in the opposite direction shows how decisive the rollout-length choice is: by removing late-position supervision, the student preserves its own succinct style while still inheriting the teacher’s reasoning strategy, leading to a more desirable outcome than simply copying the teacher.

Furthermore, we examine quantitatively how the student’s probability output aligns with the teacher’s modes. We take the top-10% highest-KL tokens after training ($n=110,894$), which reveal behavioral differences most saliently, and check how often the student’s top choice agrees with the teacher’s top-1, falls in the teacher’s top 2–5, or lies outside the top-5 (Table 4, right). We find that, indeed, ESR produces a model that is more committed to the teacher’s top 2–5 choices than to the teacher’s top-1 (47.4% in top 2–5 vs 44.6% for full-rollout; 41.9% argmax agreement vs 45.7%). At the same time, ESR’s top-1 probability is higher than full-rollout’s (0.79 vs 0.77), showing the student is also more *confident* in its own chosen token. This exactly shows that ESR steers the student to commit to a secondary mode in the teacher’s distribution.

6 Related Work

Knowledge Distillation for Language Models. Knowledge distillation [Hinton et al., 2015] transfers knowledge from a teacher to a smaller student via soft targets, and Kim and Rush [2016] extended this idea to sequence models with word-level and sequence-level objectives. For autoregressive LLMs, both the divergence and the data distribution are crucial. Gu et al. [2024] advocated reverse KL for generative LLM distillation, arguing that it avoids assigning mass to low-support teacher regions, and Agarwal et al. [2024] introduced Generalized Knowledge Distillation (GKD), which uses student-generated rollouts to obtain substantial gains over off-policy distillation on reasoning tasks. Related work has explored other divergence and sampling choices, including skew KL and adaptive off-policy schedules [Ko et al., 2024], general f -divergences [Wen et al., 2023], the mode-seeking versus mean-seeking behavior of forward and reverse KL [Wu et al., 2025], and speculative knowledge distillation with interleaved teacher-student sampling [Xu et al., 2025]. Our work builds directly on the on-policy reverse-KL setting of Gu et al. [2024] and Agarwal et al. [2024], but asks a different question: holding the divergence and rollout distribution fixed, which token positions carry useful signal?

Token-Level Importance in Distillation and Reasoning. A growing line of work suggests that not all tokens contribute equally to learning. In reasoning, Wang et al. [2025] found that only a small fraction of chain-of-thought tokens are high-entropy “forking tokens” that steer subsequent reasoning,

while Vassoyan et al. [2025] showed that uniform KL penalties can suppress exploration on critical tokens and proposed entropy-weighted KL relaxation. Related studies also identify token-level structure in planning and credit assignment, including preplan-and-anchor behavior [Li et al., 2025] and functional importance in reasoning chains [Singh and Hakkani-Tür, 2026].

In distillation specifically, several concurrent methods have explored token selection or weighting. SelecTKD uses teacher verification to mask rejected tokens [Huang et al., 2025]; AdaKD adapts token-level temperature based on training stability [Xie et al., 2026]; SE-KD disentangles selection along position, class, and sample axes and uses student-entropy filtering along the position axis [Tavor et al., 2026]; and TSDKD combines entropy-based token selection with preference ranking [Kim and Baek, 2026]. Our ablation in Section 5.1 shows that these scalar token-saliency criteria are insufficient: selecting the same number of tokens by top-KL, top-entropy, or combined entropy heuristics all underperform ESR, and most underperform even plain full-sequence training. This indicates that position is a load-bearing axis of supervision rather than a proxy for token saliency. It also reconciles our findings with Wang et al. [2025]: in on-policy distillation, the high-entropy “forking tokens” are concentrated in the early uncontaminated window, so their conclusion is consistent with ours once position is taken into account.

Concurrent work. Zhang et al. [2026] independently report that concentrating OPD supervision on response prefixes is an effective efficiency lever; in their setting — distilling a reasoning teacher into a base model that has not yet acquired reasoning behavior — prefix OPD does not surpass full-trajectory OPD, suggesting that bootstrapping reasoning from scratch still benefits from full-trajectory supervision. Our setting is complementary: starting from a math-SFT student that already reasons, on-policy distillation pushes the student *beyond* the teacher, an effect we identify and empirically attribute to the mode-seeking nature of reverse-KL; we additionally provide systematic experiments and analyses on why late-position tokens are detrimental and on the mechanism through which prefix tokens drive learning — none of which is addressed by concurrent work.

7 Conclusion

We presented Early Stopping Rollout (ESR), a one-line modification to on-policy distillation that restricts both the rollout and the reverse-KL loss to the first N response tokens. Despite its simplicity, ESR is simultaneously *better*, *more efficient*, and *stabler* than full-rollout OPD: it matches or surpasses the baseline across tasks (math, coding, function calling), regimes (LoRA, full FT), scales (1.5B–7B students, 1.7B–14B teachers), and families (Qwen, Gemma), while cutting wall-clock by $\sim 24\times$ and peak GPU memory by $\sim 2.6\times$ and eliminating the HumanEval, BFCL, and MATH-500 collapse modes observed under full-rollout training. We identify four complementary reasons for the gain: (i) *off-policy teacher drift* — late-position teacher scores are conditioned on student-specific, off-distribution prefixes and degrade monotonically with position; (ii) *stale late-position lock-in* — the right late-position target shifts as earlier tokens improve, so full-rollout training carries ballast from earlier prefixes; (iii) a *convergence cascade* that drops KL on the untrained tail by 30–40%, making explicit late supervision unnecessary; and (iv) position itself is an independent supervisory axis, since the same number of tokens selected by top-KL or top-entropy underperforms ESR and often even full-rollout. Most strikingly, ESR-trained students frequently *exceed* the teacher itself — an effect we attribute to the mode-seeking nature of reverse-KL: supervising only the planning-rich early window lets the student commit to a non-argmax but teacher-supported mode while preserving its own succinct style. This suggests that distilled students can in principle surpass their teachers when the loss interacts favorably with the divergence’s geometry, and motivates future work on which positions, divergences, and schedules best exploit this opening.

Limitations. Our experiments are conducted on small open-source models with a limited data budget; whether the ESR story holds at industrial scale (much larger students, teachers, and training corpora) remains unclear. We also do not test multi-modality or long-horizon tasks, both of which may exhibit different positional signal-quality patterns.

Broader Impacts. We present a highly effective method that enables both academic and industrial practitioners to perform on-policy distillation at substantially lower compute and memory cost — roughly $24\times$ faster wall-clock and $2.6\times$ lower peak memory in our experiments. By drastically reducing the resource floor for training competitive reasoning models, our approach helps democratize

AI training, making state-of-the-art distillation accessible to smaller labs and groups that cannot afford the compute budgets typically assumed in this line of work.

References

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*, 2024.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Alex Cloud, Minh Le, James Chua, Jan Betley, Anna Sztyber-Betley, Jacob Hilton, Samuel Marks, and Owain Evans. Subliminal learning: Language models transmit behavioral traits via hidden signals in data. *arXiv preprint arXiv:2507.14805*, 2025.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large language models. In *International Conference on Learning Representations*, 2024.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems*, 2021.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. In *NIPS Deep Learning and Representation Learning Workshop*, 2015.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Haiduo Huang, Jiangcheng Song, Yadong Zhang, and Pengju Ren. SelecTKD: Selective token-weighted knowledge distillation for LLMs. *arXiv preprint arXiv:2510.24021*, 2025.
- Minsang Kim and Seung Jun Baek. Explain in your own words: Improving reasoning via token-selective dual knowledge distillation. *arXiv preprint arXiv:2603.13260*, 2026.
- Yoon Kim and Alexander M. Rush. Sequence-level knowledge distillation. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 1317–1327, 2016.
- Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. DistiLLM: Towards streamlined distillation for large language models. In *International Conference on Machine Learning*, 2024.
- Yang Li, Zhichen Dong, Yuhan Sun, Weixun Wang, Shaopan Xiong, Yijia Luo, Jiashun Liu, Han Lu, Jiamang Wang, Wenbo Su, Bo Zheng, and Junchi Yan. Attention illuminates LLM reasoning: The preplan-and-anchor rhythm enables fine-grained policy optimization. *arXiv preprint arXiv:2510.13554*, 2025.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*, 2023.
- Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems*, 2023.
- Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48371–48392, 2025.

- 357 Qwen Team. Qwen2.5-math technical report: Toward mathematical expert model via self-
358 improvement. *arXiv preprint arXiv:2409.12122*, 2024.
- 359 Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- 360 Janvijay Singh and Dilek Hakkani-Tür. Do LLMs encode functional importance of reasoning tokens?
361 *arXiv preprint arXiv:2601.03066*, 2026.
- 362 Almog Tavor, Itay Ebenspanger, Neil Cnaan, and Mor Geva. Rethinking selective knowledge
363 distillation. *arXiv preprint arXiv:2602.01395*, 2026.
- 364 Jean Vassoyan, Nathanael Beau, and Roman Plaud. Ignore the KL penalty! boosting exploration
365 on critical tokens to enhance RL fine-tuning. In *Findings of the Association for Computational*
366 *Linguistics: NAACL 2025*, 2025.
- 367 Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
368 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
369 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 rule: High-entropy minority tokens drive
370 effective reinforcement learning for LLM reasoning. In *Advances in Neural Information Processing*
371 *Systems*, 2025.
- 372 Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f-divergence minimization for sequence-level
373 knowledge distillation. In *Proceedings of the 61st Annual Meeting of the Association for Computa-*
374 *tional Linguistics (Volume 1: Long Papers)*, pages 10817–10834, 2023.
- 375 Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. Rethinking
376 Kullback-Leibler divergence in knowledge distillation for large language models. In *Proceedings*
377 *of the International Conference on Computational Linguistics (COLING)*, 2025.
- 378 Xurong Xie, Zhucun Xue, Jiafu Wu, Jian Li, Yabiao Wang, Xiaobin Hu, Yong Liu, and Jiangn-
379 ing Zhang. LLM-oriented token-adaptive knowledge distillation. In *Proceedings of the AAAI*
380 *Conference on Artificial Intelligence*, 2026.
- 381 Wenda Xu, Rujun Han, Zifeng Wang, Long T. Le, Dhruv Madeka, Lei Li, William Yang Wang,
382 Rishabh Agarwal, Chen-Yu Lee, and Tomas Pfister. Speculative knowledge distillation: Bridging
383 the teacher-student gap through interleaved sampling. In *International Conference on Learning*
384 *Representations*, 2025.
- 385 Dongxu Zhang, Zhichao Yang, Sepehr Janghorbani, Jun Han, Andrew Ressler II, Qian Qian, Gre-
386 gory D. Lyng, Sanjit Singh Batra, and Robert E. Tillman. Fast and effective on-policy distillation
387 from reasoning prefixes. *arXiv preprint arXiv:2602.15260*, 2026.

388 A Training Efficiency

Table 5: **Training efficiency, detailed breakdown.** Per-step wall-clock time on a single A6000 (48GB). Student: Qwen2.5-Math-1.5B, LoRA, bs=16.

Teacher	Method	Gen (s)	Score (s)	Train (s)	Total (s)
Q3-1.7B	Full-seq	100–731	3–10	5–12	~280
	Ours ($N=100$)	~5	~1	~2	~8
Q3-4B	Full-seq	97–733	3–9	6–10	~210
	Ours ($N=100$)	~5	~1	~2	~8
Q3-8B	Full-seq	96–230	7–11	1–6	~170 [†]
	Ours ($N=100$)	~5	~1	~2	~8

[†] Frequent OOMs; 48GB insufficient for 8B teacher + vLLM + student.

Table 6: **Additional student-teacher pairs.** All use ESR-100, LoRA, 3,200 problems.

Student	Teacher	MATH-500 avg@4	BFCL full_acc	HumanEval pass@1
Qwen3-1.7B	— (baseline)	69.20%	59.80%	39.63%
Qwen3-1.7B	Qwen3-4B	69.20%	73.70%	44.51%
Qwen3-1.7B	Qwen3-8B	69.15%	74.00%	40.24%
Qwen3-4B	— (baseline)	77.95%	73.00%	73.17%
Qwen3-4B	Qwen3-8B	77.05%	76.80%	71.34%

B Additional Student-Teacher Pairs

C Full Experimental Results

C.1 Math Results: Per-Step Performance

Table 7 presents per-step results for the primary math experiments (LoRA, $n=1$, 3,200 problems).

Table 7: **MATH-500 per-step results.** LoRA, $n=1$, bs=16, 3,200 problems. Baseline: 50.95% avg@4.

Method	Metric	Step 50	Step 100	Step 150	Step 200
ESR-50	avg@4	62.35	66.05	66.65	64.85
	maj@4	69.40	72.00	71.00	71.20
	pass@4	77.20	79.40	81.00	79.60
ESR-100	avg@4	63.75	64.45	65.15	65.85
	maj@4	70.00	68.40	69.60	70.80
	pass@4	79.80	78.40	80.20	79.80
ESR-150	avg@4	65.35	66.65	65.30	65.75
	maj@4	66.80	67.00	66.30	67.30
	pass@4	79.00	81.00	78.20	80.00
ESR-200	avg@4	66.05	64.65	65.10	65.55
	maj@4	71.20	68.40	70.00	71.20
	pass@4	81.00	79.80	80.60	80.60

C.2 Math Results: Earlier Series ($n=16$, 200 problems)

C.3 Coding Results

C.4 Function Calling Results

C.5 Scaling Experiment Results

C.6 Full-Sequence Degradation Details

Table 12 summarises every full-rollout cell (across pairs, tasks, and regimes) for which we have a per-step trajectory on disk and observe degradation across training. For each cell we report the *peak* score (best checkpoint across training) and the *lowest* score (worst checkpoint after the peak), together with the step indices at which they occur. Bold Δ values indicate cells where Full Rollout drops below the student baseline. The Qwen-math $n=16$ row (boxed-repetition collapse) is documented in greater detail in Table 13 below.

Per-step trajectories for the seven cells in Table 12:

- **Qwen-1.5B→Q3-1.7B, LoRA, HumanEval (pass@1):** 40.2 (s50) → 31.7 → 32.3 → 32.9 → 27.4 → 28.0 → 26.8 → 26.8 (s400).
- **Qwen-1.5B→Q3-1.7B, LoRA, HumanEval+ (pass@1):** 35.4 (s50) → 27.4 → 29.9 → 29.9 → 25.0 → 26.2 → 25.0 → 25.0 (s400).

Table 8: **Complete MATH-500 results, LoRA, $n=16$, $bs=16$.** Best per configuration in **bold**.

Config	Step 50	Step 100	Step 150	Step 200
<i>avg@4</i>				
ESR-5	55.15	55.65	56.50	56.30
ESR-10	59.50	56.40	56.30	56.00
ESR-20	60.20	57.35	56.40	56.55
ESR-50	62.45	62.15	60.50	61.90
ESR-100	61.15	63.55	64.25	63.85
ESR-150	63.55	65.70	64.80	64.40
ESR-200	63.25	66.75	64.55	65.25
Progressive	61.10	62.05	61.95	62.30
Full-seq [†]	65.60	46.30	49.70	43.80
Full-seq (first-boxed)	65.55	64.95	65.55	64.95
<i>pass@4</i>				
ESR-5	74.0	74.8	76.0	73.4
ESR-10	75.8	73.2	75.4	74.2
ESR-20	77.0	76.0	73.4	75.8
ESR-50	77.0	78.4	77.6	78.6
ESR-100	77.6	79.2	80.2	80.2
ESR-150	79.6	77.4	80.0	79.4
ESR-200	79.6	80.2	79.4	79.2
Progressive	77.2	77.8	78.2	78.8
Full-seq	80.0	72.2	77.0	74.2

[†] Last-boxed extraction; degrades due to repetition.

Table 9: **Complete coding results, LoRA.** HumanEval (HE) and MBPP pass@1.

Config	Metric	s50	s100	s150	s200	s250	s300	s350	s400
ESR-50	HE	37.8	39.0	39.6	41.5	40.2	40.9	42.1	40.9
	MBPP	52.1	48.9	46.0	46.0	47.6	46.3	46.6	46.3
ESR-100	HE	37.2	39.0	42.1	37.8	39.0	37.8	37.8	38.4
	MBPP	52.1	49.7	49.2	49.2	48.4	49.2	47.9	48.9
ESR-150	HE	36.6	35.4	36.6	39.0	41.5	39.6	38.4	37.2
	MBPP	51.3	50.0	48.4	50.5	49.5	50.3	49.7	48.7
Full-seq	HE	40.2	31.7	32.3	32.9	27.4	28.0	26.8	26.8
	MBPP	52.6	48.9	49.5	48.1	47.9	47.1	48.4	47.6

Table 10: **Function calling results (BFCL), LoRA.** Name accuracy / Full accuracy / Parse rate. Best full_acc in **bold**.

Method	Best Step	Name Acc	Full Acc	Parse Rate
Baseline	—	9.70%	2.70%	24.20%
Teacher (Qwen3-1.7B)	—	75.30%	54.00%	75.30%
ESR-50	200	95.20%	57.20%	98.30%
ESR-100	100	86.20%	61.30%	91.30%
ESR-150	200	88.70%	61.50%	92.50%
ESR-200	200	80.80%	54.50%	90.20%
Full-seq	100	81.00%	58.20%	86.70%

Table 11: **Scaling results.** All use ESR-100, LoRA, 3,200 problems. Best across training steps. Baseline performance shown for reference.

Student → Teacher	Math avg@4	BFCL full_acc	HE pass@1	MBPP+ pass@1
M-1.5B baseline	50.95%	2.70%	—	—
M-1.5B → Q3-1.7B	65.85%	61.30%	42.07%	44.44%
M-1.5B → Q3-4B	68.95%	45.80%	43.29%	47.62%
M-1.5B → Q3-8B	67.85%	57.00%	41.46%	44.44%
Q3-1.7B baseline	69.20%	59.80%	39.63%	44.71%
Q3-1.7B → Q3-4B	69.20%	73.70%	44.51%	47.62%
Q3-1.7B → Q3-8B	69.15%	74.00%	40.24%	48.15%
Q3-4B baseline	77.95%	73.00%	73.17%	55.56%
Q3-4B → Q3-8B	77.05%	76.80%	71.34%	56.35%

Table 12: **Full-rollout collapses across training (peak → lowest).** Each row is a single Full Rollout experiment; we report the highest and lowest checkpoint scores across all saved steps. “Δ” is lowest – peak. Cells with [†] end the run *below the student baseline*. Trajectories are taken from the per-step eval logs on the training cluster (docs/main_results.md, docs/scaling_results.md). ESR on the same cells (not shown here) stays within ±1pp of its peak across training; see Tables 7 and 9.

Pair	Regime	Task	Metric	Peak (step)	Lowest (step)	Δ
Qwen-1.5B → Q3-1.7B	LoRA	MATH-500 ($n=16$)	avg@4	65.60 (s50)	43.80 (s200) [†]	−21.80
Qwen-1.5B → Q3-4B	LoRA	MATH-500	avg@4	67.45 (s50)	54.45 (s100)	−13.00
Qwen-1.5B → Q3-4B	LoRA	MATH-500	maj@4	73.20 (s50)	63.40 (s200)	−9.80
Qwen-1.5B → Q3-1.7B	LoRA	HumanEval	pass@1	40.2 (s50)	26.8 (s350)	−13.4
Qwen-1.5B → Q3-1.7B	LoRA	HumanEval+	pass@1	35.4 (s50)	25.0 (s250)	−10.4
Qwen-1.5B → Q3-1.7B	FullIFT	HumanEval	pass@1	36.6 (s150)	30.5 (s300)	−6.1
Qwen-1.5B → Q3-1.7B	FullIFT	HumanEval+	pass@1	31.7 (s150)	26.2 (s250)	−5.5
Gemma-2-2B → Gemma-3-4B	LoRA	HumanEval	pass@1	22.0 (s50)	11.6 (s150) [†]	−10.4

[†] Lowest score is below the student baseline (Qwen-Math-1.5B MATH-500: 50.95%; Gemma-2-2B HumanEval: 23.78%).

- 409 • **Qwen-1.5B→Q3-1.7B, FullIFT, HumanEval (pass@1):** 32.3 → 32.3 → 36.6 (s150) →
410 36.0 → 31.1 → 30.5 → 31.7 → 31.1 (s400).
- 411 • **Qwen-1.5B→Q3-1.7B, FullIFT, HumanEval+ (pass@1):** 27.4 → 28.7 → 31.7 (s150) →
412 30.5 → 26.2 → 26.2 → 26.8 → 26.2 (s400).
- 413 • **Gemma-2-2B→Gemma-3-4B, LoRA, HumanEval (pass@1):** 22.0 (s50) → 14.6 → 11.6
414 → 11.6 (s200). Falls below the 23.78 student baseline at s100 and stays there.
- 415 • **Qwen-1.5B→Q3-4B, LoRA, MATH-500 avg@4:** 67.45 (s50) → 54.45 → 58.85 → 55.05
416 (s200) (boxed-repetition pattern; ESR-100 on the same pair climbs to 68.95).
- 417 • **Qwen-1.5B→Q3-1.7B, LoRA, MATH-500 ($n=16$) avg@4:** 65.60 (s50) → 46.30 → 49.70
418 → 43.80 (s200). Boxed-repetition collapse; statistics in Table 13.

Table 13: **Boxed repetition progression** in full-sequence LoRA distillation on MATH-500 ($n=16$, bs=16). Companion to the corresponding row of Table 12.

Statistic	Step 50	Step 100	Step 150	Step 200
Multi-boxed responses	2.9%	91.5%	80.5%	89.1%
Avg \boxed{} per response	1.0	88.1	54.1	58.0
Repetitive responses	5.9%	78.8%	58.3%	77.5%
Avg response length (chars)	1,388	4,023	3,585	4,558
Official avg@4	65.60%	46.30%	49.70%	43.80%
Relaxed avg@4 (GT in text)	74.60%	70.90%	71.90%	70.30%

D KL-Profile Diagnostic for Choosing N

A practical question for a new student–teacher pair is: what value of N should I use? Our position-sweep (Section 4.3) shows a wide plateau on Qwen math ($N \in [50, 200]$, avg@4 within $\sim 1\text{pp}$); the same plateau appears on Gemma math ($N \approx 100$). These plateau locations coincide with the position at which the cumulative per-position KL $\sum_{k \leq N} \mathbb{E}[\text{KL}_k]$ reaches $\sim 45\%$ of the full-trajectory total, measured on a small ($\sim 1\text{k}$ trajectory) pre-training rollout. This suggests a cheap rule of thumb: run a 1k-trajectory per-position KL scan before training, and pick N at the 45% cumulative-KL point. We verify this on two cells (Qwen, Gemma math) and leave a full six-cell verification ($\{\text{Qwen, Gemma}\} \times \{\text{math, coding, funcall}\}$) as a short follow-up (TBD: KL-profile scan); the diagnostic is reported here as a practical recipe rather than a claim.

E Token Classification Methodology

We classify each token into six categories based on string matching:

1. **planning**: Reasoning keywords (“To”, “Let”, “First”, “Step”, “We”, “Given”, “Therefore”, “Thus”, “Since”).
2. **structural**: Punctuation, whitespace, formatting tokens.
3. **math_number**: Digits (0–9).
4. **math_operator**: Arithmetic operators (+, −, ×, /, =).
5. **math_latex**: LaTeX delimiters ($\backslash$ (, $\backslash$).
6. **continuation**: All others.

Table 14: Mean KL by token category and position range.

Category	0–4	5–19	20–49	50–99	100–199	200–499
planning	4.50	0.79	1.49	1.66	1.49	2.37
structural	3.26	1.46	1.60	0.93	0.60	0.86
math_number	1.49	0.60	0.74	0.28	0.17	0.13
math_operator	7.30	1.84	0.81	0.37	0.21	0.14
math_latex	8.84	9.23	6.50	4.95	2.97	1.87
continuation	1.94	1.12	1.19	0.89	0.70	0.48

F Assets and Licenses

Table 16 lists the models and datasets used in this paper, with their providers and licenses. All assets are used in accordance with their respective terms of use.

Table 15: **Top 20 highest-KL tokens** (minimum 50 occurrences across 10,000 trajectories).

Rank	Token	Category	Count	Mean KL
1	“Solution”	planning	152	21.93
2	“Analysis”	continuation	125	16.49
3	\[	math_latex	7,152	13.21
4	“examines”	continuation	74	11.51
5	“He”	continuation	150	10.80
6	\(	math_latex	21,243	10.30
7	“First”	planning	1,706	9.98
8	“tests”	continuation	52	8.71
9	\\	math_latex	82	8.68
10	“There”	continuation	201	8.28
11	“Therefore”	planning	4,913	7.95
12	“Identify”	continuation	1,345	7.78
13	“To”	planning	8,806	6.90
14	“When”	continuation	89	6.62
15	“This”	continuation	621	6.34
16	“Thus”	planning	2,547	6.24
17	“First” (space)	planning	259	6.10
18	“To” (space)	planning	1,174	5.37
19	“The”	planning	2,626	5.25
20	“Next”	planning	1,753	5.03

Table 16: **Assets used in this paper.** All licenses verified at time of submission.

Asset	Type	Provider	License
Qwen2.5-Math-1.5B	Model	Alibaba	Apache 2.0
Qwen2.5-Math-7B	Model	Alibaba	Apache 2.0
Qwen3-1.7B / 4B / 8B / 14B	Model	Alibaba	Apache 2.0
Gemma-2-2B	Model	Google	Gemma Terms of Use
Gemma-3-4B	Model	Google	Gemma Terms of Use
NuminaMath	Dataset	Numina	Apache 2.0
CodeUltraFeedback	Dataset	Coseal	MIT
glaive-function-calling-v2	Dataset	Glaive AI	Apache 2.0
MATH-500	Benchmark	Hendrycks et al.	MIT
HumanEval / HumanEval+	Benchmark	OpenAI / EvalPlus	MIT
BFCL	Benchmark	UC Berkeley	Apache 2.0

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and §1 state that ESR matches or exceeds full-rollout distillation while being $20\text{--}35\times$ faster with $\sim 4\times$ less memory; each claim maps to Table 1, §4.3, §4.4, and §5.1.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: A *Limitations* paragraph at the end of §7 discusses the small open-source-model regime, limited data budget, and untested industrial scale.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is empirical; no formal theorems or proofs.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §4 lists student/teacher pairs, datasets, all hyperparameters, the loss, and evaluation protocols; Appendices A and C give per-step timing and results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: We will release the training code and LoRA adapters publicly after acceptance; the methodology in §4 is detailed enough to reproduce the main results in the meantime.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: All training and evaluation settings are stated in the Setup paragraph of §4, with per-step results in Appendix C.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Reported numbers are single-seed best-step results; robustness comes from breadth across families, tasks, and teacher scales rather than per-cell variance, and this is acknowledged in §7.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix A (Table 5) reports per-step wall-clock on a single A6000 (48GB) for each teacher scale, broken down by phase.

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We use open-weight models and public benchmarks, with no human subjects, private data, or scraped content.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: A *Broader Impacts* paragraph at the end of §7 discusses both the positive impact of cheaper distillation research and the dual-use risk of lower-cost capability extraction from proprietary teachers.

11. Safeguards

507 Question: Does the paper describe safeguards that have been put in place for responsible
508 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
509 image generators, or scraped datasets)?

510 Answer: [N/A]

511 Justification: No new pretrained base models or scraped datasets are released; trained
512 students are LoRA adapters on top of public bases.

513 **12. Licenses for existing assets**

514 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
515 the paper, properly credited and are the license and terms of use explicitly mentioned and
516 properly respected?

517 Answer: [Yes]

518 Justification: All models and datasets are cited in §4 and a consolidated license table is
519 provided in Appendix F.

520 **13. New assets**

521 Question: Are new assets introduced in the paper well documented and is the documentation
522 provided alongside the assets?

523 Answer: [N/A]

524 Justification: No new assets are released at submission; LoRA adapters and code will be
525 documented and released after acceptance.

526 **14. Crowdsourcing and research with human subjects**

527 Question: For crowdsourcing experiments and research with human subjects, does the paper
528 include the full text of instructions given to participants and screenshots, if applicable, as
529 well as details about compensation (if any)?

530 Answer: [N/A]

531 Justification: No crowdsourcing or human subjects.

532 **15. Institutional review board (IRB) approvals or equivalent for research with human**
533 **subjects**

534 Question: Does the paper describe potential risks incurred by study participants, whether
535 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
536 approvals (or an equivalent approval/review based on the requirements of your country or
537 institution) were obtained?

538 Answer: [N/A]

539 Justification: No human subjects research.

540 **16. Declaration of LLM usage**

541 Question: Does the paper describe the usage of LLMs if it is an important, original, or
542 non-standard component of the core methods in this research? Note that if the LLM is used
543 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
544 scientific rigor, or originality of the research, declaration is not required.

545 Answer: [N/A]

546 Justification: LLMs are the object of study in this work and do not require declaration. We
547 additionally used LLMs as coding and writing assistants, but all methodological choices,
548 experimental design, analyses, and conclusions were author-led; the LLMs did not contribute
549 original methodology.