
Less is More: Early Stopping Rollout for On-Policy Distillation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 On-policy distillation has recently emerged as a promising alternative to standard
2 sequence-level imitation, training a student by scoring its own rollouts with a
3 teacher model. However, one important limitation is observed in this paradigm:
4 at each position, the teacher scores the student’s next token conditioned on the
5 student’s own previous outputs. Early tokens are therefore judged mainly in the
6 context of the original problem and reflect how the student approaches the problem,
7 whereas later tokens are constrained by increasingly student-specific context and
8 provide much weaker supervision. Based on this intuition, we propose Early
9 Stopping Rollout (ESR), a simple yet effective distillation strategy that restricts
10 both rollout generation and reverse-KL to the first N response tokens. This position-
11 focused objective emphasizes the tokens that encode planning and reasoning, rather
12 than spending supervision budget on late tokens whose informational value rapidly
13 degrades. Empirical experiments have proved that ESR shows both surpassing
14 performance with high efficiency under various models and training methods. It
15 is also proved that the teacher collapses in sufficiently long rollouts suggesting
16 that late-position supervision contributes little useful learning signal. We further
17 investigate the mechanism behind this effect and show that position is a distinct
18 and crucial dimension of supervision, not merely a proxy for token uncertainty or
19 difficulty. It establishes position-aware supervision as an important direction for
20 understanding and advancing OPD in future studies.

21 1 Introduction

22 On-policy distillation (OPD) has emerged as a dominant paradigm for model distillation in industrial
23 practice. The student generates its own rollouts τ , which are then scored by the teacher: at each token,
24 the teacher’s probability $\pi_{teacher}(\tau_{student} \mid x_{prompt})$ serves as the soft target for the student [Agarwal
25 et al., 2024, Gu et al., 2024]. Viewed through an RL lens, this is equivalent to using the teacher as a
26 dense, token-level reward model that judges the student’s own behavior on a given prompt. The final
27 objective is then a simple *uniform average* of the per-token reverse-KL across the trajectory.

28 This uniform averaging implicitly assumes that the teacher’s signal is equally reliable at every position
29 in the trajectory. We argue that this assumption is unfounded, and that it breaks down precisely
30 because of the autoregressive structure of LLMs. At the first token, the teacher’s score is conditioned
31 only on the prompt: $\pi_{teacher}(\tau_{student}^{t=1} \mid x_{prompt})$ - indeed what we expect the teacher to score on. At
32 a later position m , however, it becomes conditioned also on the student’s own previously generated
33 tokens: $\pi_{teacher}(\tau_{student}^{t=m} \mid x_{prompt}, \tau_{student}^{t=1:m})$. This conditioning context is off-policy to the teacher
34 model and may even be out-of-distribution (OOD), and grows more so with position.

35 To probe how severe this drift is, we run a preliminary experiment that measures how the teacher
36 behaves when continuing from an early-stopped student rollout. As shown in Figure 1, the teacher’s

quality collapses toward the student’s within roughly 300 tokens. If the teacher’s late-position signal is unreliable, the natural question is whether early tokens alone are systematically more informative — and whether training on them is sufficient.

Motivated by this, we propose Early Stopping Rollout(ESR): restrict the student rollout to its first N tokens (e.g., 50 or 100) and compute the distillation loss only on this early window. The change is a single line in any on-policy distillation loop. Despite its simplicity, ESRconsistently outperforms full-rollout OPD across tasks (math, code, function calling), training regimes (LoRA, full fine-tuning), model scales (students 1.5B–7B, teachers 1.7B–14B), and model families (Qwen, Gemma), while reducing wall-clock cost by up to $24\times$ and peak training memory by up to $4\times$. We further observe that full-rollout OPD frequently destabilizes and collapses during training, whereas ESRdoes not collapse in any of our settings. Perhaps most surprisingly, although the student is normally upper-bounded by the teacher’s performance, ESR-trained students often *exceed* the teacher.

Such a small change yields such large gains probe us to investigate deeper. Firstly, we want to further validate whether the effectiveness of earlier tokens is intrinsically caused by better KL and entropy signals, because we also find that early tokens generally have larger KL and entropy. We test selectors that pick tokens by high KL, high entropy or combined regardless of position; all underperform ESR, with most underperforming full-rollout. Therefore, although we don’t exclude KL and entropy as auxillary factors, we confirm that the position is a an independent factor rather a saliency proxy. Furthermore, we identified an additional important mechanism *Convergence Cascade*: after training on the early window, KL divergence on the *untrained* late tokens also drops by 30–40%. Therefore, the KL divergence does not have to see late positions to repair them.

Lastly, we investigated into why ESRsometimes even *exceeds* the teacher — a stronger phenomenon we observe on most of our main experiment Table 1. We trace this to the mode-seeking nature of reverse-KL: the student commits to a subset of supported teacher modes instead of chasing all of them. Our ESRloss helps it steer toward a non-dominant but correct mode rather than the teacher’s argmax. Empirically, we find ESR-trained students commit harder to their own choices than full-rollout students yet agree *less* with the teacher’s argmax, and produce trajectories $\sim 3\times$ shorter than full-rollout and $\sim 2\times$ shorter than the base student (§5.3). It learns the teacher’s strategy but preserves its succinct style, leading to better compound behaviors. This finding showcases a potential path of superceding the teacher model in distillation that worth further investigations.

Our contributions are summarized below:

1. **Early Stopping Rollout(ESR) outperforms full-rollout on-policy distillation.** A one-line change—restricting the rollout length to the first N response tokens—beats full-sequence distillation across tasks, model families, scales, and training regimes, while being dramatically more efficient and far more stable to train. It can even exceed the teacher’s own performance (Sections 3–4.3).
2. **A systematic investigation of why it works.** First, through a token-selection ablation, we show that the gain is not “high KL” or “high entropy” in disguise: position itself is a load-bearing axis. We also identify convergence cascading to be another important mechanism. Furthermore, we explain why ESRoften surpasses the teacher: by exploiting the mode-seeking behavior of reverse-KL training, ESRpushes the student toward a better, non-dominant sub-mode (Section 5).

2 Why Late-Position Supervision Is Unreliable

In on-policy distillation, the teacher model scores the student’s rollout at each position t via $\pi_t(\cdot \mid x, y_{<t}^{\text{student}})$ and averages these scores uniformly across the rollout. The conditioning context $y_{<t}^{\text{student}}$, however, is whatever this particular student happened to write at this particular moment of training. As t grows it drifts away from any distribution the teacher itself was trained on—off-policy to the teacher and, for long enough or atypical enough rollouts, plausibly out-of-distribution. We term this problem as “Off-policy Teacher Drift”.

To verify the intuition, we feed the teacher a k -token student-generated rollout on MATH-500 and letting it continue autoregressively, the teacher’s avg@4 ($n=4$, $t=0.7$) falls from its 65.30% unconditional baseline to 62.70% at $k=100$ and 51.75% at $k=300$ —essentially student-baseline

(Figure 1). Late-position teacher scores are not independent assessments of the prompt; it is a guess at how to continue a trajectory the student has already committed to.

The problem compounds during training because the student’s policy is moving. Early on, the student’s rollout is far from the teacher; late-position targets are fitted under a strongly off-policy prefix. As training progresses and earlier tokens approach the teacher’s distribution, the right late-position behavior changes too—continuations that were sensible above an off-policy prefix are no longer sensible above an on-policy one. But the late-position habits learned earlier do not cleanly erase: gradient descent gives no guarantee of unlearning, and in LoRA’s low-rank setting it updates make redirection compete for capacity that may already be spent. Full-rollout training thus chases a moving late-position target while carrying ballast from earlier ones. We term this problem as “Stale Late-position Lock-in”. This insight expects a training stability issues for full-rollout OPD – a prediction we later verify in §4.3

3 Our Method: Early Stopping Rollout (ESR) On Policy Distillation

3.1 Formulation

Let π_s denote the student and π_t the teacher. In standard on-policy reverse-KL distillation, the student generates a response $\mathbf{y} = (y_1, \dots, y_T)$ conditioned on prompt x , and the loss is

$$\mathcal{L}_{\text{full}} = \mathbb{E}_{\mathbf{y} \sim \pi_s(\cdot|x)} \left[\sum_{t=1}^T \text{KL}(\pi_s(\cdot | x, \mathbf{y}_{<t}) \parallel \pi_t(\cdot | x, \mathbf{y}_{<t})) \right]. \quad (1)$$

ESR (position cutoff N , with $N \ll T$ in practice) truncates the student rollout to its first N tokens, and the loss is computed over exactly those tokens:

$$\mathcal{L}_{\text{ESR}}(N) = \mathbb{E}_{\mathbf{y} \sim \pi_s(\cdot|x), |\mathbf{y}| \leq N} \left[\sum_{t=1}^{\min(N, |\mathbf{y}|)} \text{KL}(\pi_s(\cdot | x, \mathbf{y}_{<t}) \parallel \pi_t(\cdot | x, \mathbf{y}_{<t})) \right]. \quad (2)$$

If the student emits EOS before position N , the rollout terminates naturally. Everything else—generation temperature, LoRA target modules, optimizer, scorer—is unchanged from the standard on-policy KD loop.

4 Analysis on the Advancement of ESR

4.1 Setup

Models. We evaluate across multiple student-teacher pairs. Our primary student is Qwen2.5-Math-1.5B [Qwen Team, 2024], paired with teachers of increasing scale: Qwen3-1.7B, Qwen3-4B, and Qwen3-8B [Qwen Team, 2025]. We also tested Qwen2.5-Math-7B with Qwen3-14B to test student model size generality. To test cross-family generality, we also evaluate with gemma-2-2b-it (student) and gemma-3-4b-it (teacher).

Training. All experiments use LoRA [Hu et al., 2022] ($r=32$, $\alpha=64$, all linear layers, $\text{lr} = 5 \times 10^{-5}$) with reverse KL divergence loss. Temperature 0.7 for on-policy generation. Each training step processes a batch of 16 problems with 1 rollout per problem ($n_{\text{samples}}=1$, batch size 16). We train for 200 steps on all tasks, and saving checkpoints every 50 steps. Training data are drawn from NuminaMath, CodeUltraFeedback, and glaive-function-calling-v2. Our method uses $N=100$ unless otherwise specified.

Evaluation. MATH-500 [Hendrycks et al., 2021, Lightman et al., 2023] with $n=4$ samples at temperature 0.7 (reporting avg@4), HumanEval/HumanEval+ [Chen et al., 2021, Austin et al., 2021, Liu et al., 2023] at temperature 0.0 (reporting pass@1), BFCL [Patil et al., 2025] with 600 problems reporting full accuracy: correct function name and arguments.

4.2 Overall performance

Table 1 presents the main results across tasks and model families. Table 2 presents ablation results with full finetuning. Table 2 shows model size scaling results. We report the best performance

Table 1: **Main results: ESR dominates Full Rollout across tasks and families.** Best score across training steps; metrics defined in §4 setup. ★: ESR surpasses the teacher reference. **Full Rollout** numbers mark the best step before degradation; ↓ flags cells where Full Rollout further decays with training (more details in Appendix C).

Pair	Method	MATH-500		Coding		BFCL
		avg@4	pass@4	HE	HE+	full_acc
Qwen2.5-Math 1.5B → Qwen3 1.7B	Student baseline	50.95	72.80	31.10	26.20	2.70
	Teacher reference	65.30	77.00	39.60	36.00	54.00
	Full Rollout	62.35↓	74.60↓	40.20↓	35.40↓	58.20★
	ESR	65.85★	79.80★	42.10★	36.00	61.30★
Gemma-2 2B → Gemma-3 4B	Student baseline	13.45	28.20	23.78	20.10	73.17
	Teacher reference	66.60	74.80	20.70	20.10	72.83
	Full Rollout	22.95	31.40	22.00↓	17.70↓	76.83★
	ESR	27.20	39.40	28.70★	24.40★	79.00★

Table 2: **FullFT and LoRA scaling on MATH-500.** (Left) FullFT for Qwen and Gemma student/teacher pairs. (Right) LoRA scaling across teacher and student sizes; ↓ marks the **Full Rollout** peak before degradation. Both: best across training steps; **bold**: ESR beats **Full Rollout**.

Pair	Method	avg@4	pass@4
Qwen2.5-Math 1.5B	student / teacher	50.95 / 65.30	72.80 / 77.00
→ Qwen3 1.7B	Full Rollout	58.20	75.40
	ESR	56.20	73.80
Gemma-2 2B → Gemma-3 4B	student / teacher	13.45 / 66.60	28.20 / 74.80
	Full Rollout	13.90	25.00
	ESR	26.65	40.40

Pair	Method	avg@4	pass@4
Qwen2.5-Math 1.5B	student / teacher	50.95 / 77.95	72.80 / 86.40
→ Qwen3 4B	Full Rollout	67.45↓	80.60
	ESR	68.95	81.00
Qwen2.5-Math 1.5B	student / teacher	50.95 / 77.70	72.80 / 87.80
→ Qwen3 8B	Full Rollout	66.75	81.60
	ESR	67.85	82.20
Qwen2.5-Math 7B	student / teacher	53.60 / 76.15	75.00 / 83.20
→ Qwen3 14B	Full Rollout	68.85↓	80.00
	ESR	68.95	81.20

130 achieved across all training steps for each configuration. But notice that full rollout distillation
131 frequently collapse to much lower performance as it trains. We denote the sessions that collapses
132 with a down arrows.

133 **ESRbeats full rollout across model families and tasks, and frequently surpasses teacher per-**
134 **formance.** Across the cells of Table 1, ESR beats full rollout on every cell, and in six cells ESR
135 *surpasses the teacher* outright (Qwen math avg@4, pass@4; Qwen HumanEval and HumanEval+;
136 Qwen BFCL full_acc; Gemma coding on all four columns; Gemma BFCL full_acc). The Gemma
137 coding results have the biggest gap: full rollout drives HumanEval pass@1 to 22.00, *below* the 23.78
138 student baseline and worse than the 20.70 teacher, while ESR recovers 31.10. Qwen-math full rollout
139 is marked ↓ because its best number (62.35 avg@4) is a non-stable best step: under $n=16$ it collapses
140 to 46.3% by step 100 (Appendix C.6).

141 **ESRmatches or beats full rollout across full fine-tuning and LoRA settings.** Table 2 reports
142 FullFT on MATH-500 for both Qwen and Gemma. On Qwen **Full Rollout** 58.20 is slightly better
143 than ours ESR 56.20 avg@4, but the gap is close. For Gemma ESR still dominates **Full Rollout**
144 by +12.75pp avg@4 and +15.40pp pass@4. ESR is therefore the safer choice in both parameter
145 regimes.

146 **ESRscales across model sizes.** The advantage persists as both teacher and student grow. Table 2
147 reports the results scaling both the teacher and the student. With the same student (Qwen2.5-Math-
148 1.5B), we tried Qwen3-4B and 8B as the teacher. With Qwen3-4B, **Full Rollout** peaks at (67.45) and
149 then *collapses* to 55.05 (Appendix C.6), while ESR stays stable and climbs to 68.95. Against the 8B
150 teacher, ESR still edges **Full Rollout** (67.85 vs 66.75). The pattern persists at larger scale along with
151 student model too: with a Qwen2.5-Math-7B student and Qwen3-14B teacher, **Full Rollout** again
152 peaks at step 50 (68.85) and drifts to 62.40, while ESR holds stable across all checkpoints.

Figure 1: **(Left) Off-policy teacher drift.** The teacher loses accuracy quickly as the student-generated rollout grows over a few hundred tokens. MATH-500, avg@4 ($n=4$, $t=0.7$). Teacher = Qwen3-1.7B; student = Qwen2.5-Math-1.5B. After ~ 300 student tokens the teacher has effectively been dragged down to student-baseline performance. **(Right) Rollout length N sweep on MATH-500.** LoRA, Qwen2.5-Math-1.5B $\rightarrow$ Qwen3-1.7B; best avg@4 across training steps. Full rollout and the undistilled baseline are shown as horizontal references. Performance saturates for $N \in [50, 200]$ and all beat full-seq.




Table 3: **Training efficiency.** Single A6000 (48 GB), bs=16, student teacher Qwen3-1.7B. We report the average running time and GPU memory usage across student model generation, training and teacher scoring phases. Note that the real time usage can be larger if there isn’t enough GPU to hold the student and teacher models together and requires model loading and unloading.

Phase	Per-step wall time			Peak GPU memory		
	ESR (100 tok.)	Full rollout	Speedup	ESR (100 tok.)	Full rollout	Savings
Teacher scoring	1 s	7 s	7×	7.3 GB	14.9 GB	2.0×
Student generation	5 s	180 s	36×	7.2 GB	8.9 GB	1.2×
Student training	2 s	7 s	3.5×	9.6 GB	39.5 GB	4.1×
Total (single step)	8 s	194 s	24×	24.1 GB	63.3 GB	2.6×
Total (200 steps)	27 min	10.8 hr	24×			

4.3 Stability of ESR

ESR is more robust than full rollout in training. Across our (pair $\times$ task $\times$ regime) matrix, full-rollout distillation degrades with training very frequently; ESR degrades nowhere. Appendix Table 12 reports each collapsing cell as a (peak $\rightarrow$ lowest) pair with the corresponding step indices. For example, Qwen-LoRA coding falls from 40.2 to 26.8; Qwen-LoRA math collapses from 65.60 at s50 to 43.80 ($\Delta = -21.80$ pp, ending below the 50.95% student baseline); Gemma-LoRA collapse from 22.0 to 11.6 ($\Delta = -10.4$ pp). Three of seven full-rollout collapses end below the student baseline. ESR generally hold a fluctuation within ± 3 pp of its peak across training (Tables 7, 9).

Early Stopping Rollouts not sensitive to the choice of N . A natural question naturally occurs – how to choose where to stop? Is it sensitive? We conducted a set of sweeping experiments in Figure 1 sweeps N on MATH-500 and reveals a robust region: it reaches just as good performance starting from $N=50$ and remains stable to $N=200$. The method is not sensitive to the exact choice of N within a certain region.

4.4 Efficiency of ESR

Table 3 shows that ESR achieves a 24 $\times$ wall-clock speedup and reduces peak training memory by $\sim 4\times$. The dominant cost in full-seq is autoregressive generation (180 s/step for sequences averaging ~ 1000 tokens); ESR generates only $N=100$ tokens (5 s/step). Note that with ESR, all the student and teacher models can be put in one A6000 GPU comfortably. In our own practice, it saves a further big time overhead of model loading and unloading that we do not report here.

Figure 2: **(Left) Case study: early tokens encode strategy; later tokens mostly execute.** A representative MATH-500 trajectory. The first 100 tokens establish the geometric setup, name the unknown, and identify the key relationship (altitude bisects the leg)—these are the choices that determine whether the rollout will succeed. The last 100 tokens carry out an algebraic computation that any solver can finish once the strategy is fixed. ESR supervises the first window; the second follows for free. **(Right) Token selection strategies on MATH-500.** Qwen2.5-Math-1.5B $\rightarrow$ Qwen3-1.7B, LoRA, $n_{\text{samples}}=1$, best avg@4 across training steps. All selectors operate on full-length rollouts; ESR is the only one that additionally truncates the rollout. Baseline (no distillation) 50.95%; full-seq 62.35%. H_s : student entropy; H_t : teacher entropy; RKL: reverse KL.

Problem. In an isosceles right triangle, the altitude to the hypotenuse has length $4\sqrt{2}$. What is the area?

First 100 tokens (strategy). “Let’s start by understanding the problem. In an isosceles right triangle, the two legs are equal, and the hypotenuse is $\sqrt{2}$ times the leg. If we denote each leg by a , then hypotenuse $= a\sqrt{2}$. The altitude to the hypotenuse is half the leg, because the altitude bisects the...”

Last 100 tokens (execution). “...altitude from the right-angle vertex to the hypotenuse. This altitude has length $h = \frac{a \cdot a}{a\sqrt{2}} = \frac{a^2}{a\sqrt{2}} = \frac{a}{\sqrt{2}} = \frac{a\sqrt{2}}{2}$. Setting this equal to $4\sqrt{2}$: $\frac{a\sqrt{2}}{2} = 4\sqrt{2}$, so $a\sqrt{2} = 8\sqrt{2}$, giving $a = 8$. The area is $\frac{1}{2} \times 8 \times 8 = \boxed{32}$.”

Selection method	k	Best avg@4
ESR-100 (ours)	100	65.85
Full-seq (all tokens)	—	62.35
Top-RKL	100	53.35
Top- H_t (teacher entropy)	100	63.30
Top- H_s (student entropy)	100	62.70
Top- H_s	200	53.20
RKL $\cdot H_s$	100	56.90
$H_t \cdot H_s$ (product)	100	55.35
RKL $\cdot H_t \cdot H_s$ (triple product)	100	57.90

5 Additional Reasons Why Does ESR Work?

The surprising effectiveness of Early Stopping Rollout probes us to investigate deeper beyond what we have identified Section 2. We structure this section around three questions: (§5.1) Is the early-position effect induced by high-KL or high-entropy tokens? (§5.2) If not, what additional structural reasons make the early window privileged? (§5.3) Why does ESR sometimes *exceed* the teacher?

5.1 What Does The Effectiveness of ESR Stem From?

In Section 2, we identified the “Off-policy Teacher Drift” and “Stale Late-position Lock-in” problem that can cause the late-position tokens to provide degrading signals. During the experiment analysis, we find that, as shown in Figure 3, early positions simultaneously have high KL divergence between student and teacher model, and high token entropy from both student and teacher models. This finding probes us to wonder if the effectiveness is intrinsically induced by the KL and entropy. To control these factors, we conduct a series of ablation experiments: pick the same amount (100) of tokens based on the highest KL divergence, highest student/teacher entropy or with them in combination with ESR, regardless of position (Figure 2). If the effectiveness is indeed induced by the KL or entropy, then they should reproduce the same or better results.


To our surprise, all underperform ESR, and most of them also much underperform the full rollout results. Teacher or student entropy based selection can match the full sequence by them alone, but their combination falls short significantly. What’s also interesting is that KL divergence measure, the direct calculation of the loss magnitude, barely works. It only improves the baseline (50.95%) for about 3 percent. And more surprisingly, we find that the largest 100 tokens of KL occupies around 93% of the entire trajectory loss. This shows that the tokens that has larger signals are not necessarily the ones that have effective signals.

Therefore, although we still don’t exclude KL and entropy as potential mediator factor, we exclude them to be the sole factors that causes early tokens to be special. Position, therefore, should be considered an independent token selection dimension for the future.

5.2 The Convergence Cascade Effect of ESR

We find an additional important mechanism that assists in ESR’s effectiveness. Even we only trained the first N tokens with ESR, training, per-position KL divergence beyond $[0, N]$ region still drops by 30–40% (Figure 3). This shows that the student can pick up the teacher’s “global mindset” even with just the beginning tokens. One reason that we suspect is that the beginning tokens often consist of

Figure 3: **(Left)** Early tokens are high on student entropy, teacher entropy, and KL divergence simultaneously. **(Right)** **The convergence cascade.** Per-position KL between distilled student and teacher, before vs. after ESRtraining. Yellow band: positions $[0, 100]$ that actually receive training loss. Blue band: the KL gap closed by training. Positions 100+—which see no direct training signal—drop to the same KL as the trained region, confirming that alignment on the early window cascades through the autoregressive rollout.



problem framing and strategic planning content. Therefore once the student picks up how to frame problems and plan the strategy, the later content naturally follows.

Moreover, recently Cloud et al. [2025] shows student models may be able to learn the teacher’s deep internal preference even with random numbers generated by the teacher, called “subliminal learning”. Therefore, the early tokens may suffice to inject the subliminal mindset to the student.

5.3 Why Does ESR Even Exceed the Teacher?

ESR exceeds the teacher in most of the main experiments¹. Even the full rollout model slightly exceeds the teacher in function calling experiments, although weaker than our method. This shows that student has the potential exceed the teacher even in normal OPD, and our method amplifies it. Why is so? Isn’t teacher supposed to be the upper bound?

We propose the reason lies in the mechanism of reverse KL $KL(\pi_s \parallel \pi_t)$, which has a mode-seeking behavior: it penalizes the student for putting mass on tokens the teacher does not support, but not for concentrating mass on a single supported token. Therefore, the student has the possibility to land on a submode of the teacher that is actually better. We visualize this mechanism in figure 4.

To verify it is the case, we first scan through the behaviors difference of the distilled models. Surprisingly, models trained with our method produces 2x to 3x shorter sequences than not only the teacher and the full rollout versions, and even shorter than itself. Note that the teacher model is a more verbose model than the student. So learning from such teacher generally leads to longer sequences. The full rollout version is indeed the case. It produces sequence even much longer than the teacher. The juxtaposing direction learned from the same teacher caused just by the rollout length shows how important this choice leads to. Without the impact of the late-position tokens, the model preserves its original succinct style and learns the teacher’s strategy altogether, leading to a more desirable outcome than simply copying the teacher.

Furthermore, we how student’s probability output mode aligns with the teacher quantitatively. We take out the top 10 % of highest KL tokens after the training that can tell the model’s behavior difference most saliently. We check the frequency when the student’s top choice agrees with the teacher’s top choice, top 2-5 choices, and beyond. We find that, indeed, our method produces a model that is more committed to top 2-5 choices of the teacher rather than the first choice. And the top 1 choice average probability of our method is higher than that of the full rollout, showing our method leads the model to be more confident in its choice. This exactly shows that ESRsteers the model to commit to a secondary mode in the teacher model.

Behavioral signature. The hypothesis predicts that ESR-trained students should (a) produce shorter, less-hedging rollouts than the teacher, and (b) commit hard, but on *non-argmax* teacher modes. Table 4 confirms both. ESR-100 rollouts are shorter than the base student, the teacher, and full-seq across the entire length distribution — even the top-decile (~ 800 tokens) is below the base student’s median (~ 860). On the top-10% highest-KL positions ($n=110,894$), ESR raises student

Figure 4: **Mode-seeking on the planning region.** Schematic of reverse-KL behaviour at a multi-modal planning position. The teacher supports two modes (e.g. a verbose plan and a concise correct plan); reverse KL penalises student mass outside the support but not concentration within it. A short early-window loss can collapse the student onto the better supported mode, allowing the student to exceed the teacher’s average behaviour. Full-sequence training reverts the student toward the averaged teacher across late positions, undoing the concentration.




Table 4: **Behavioral Modes Analysis.** *Left:* Response length distribution on MATH-500. ESRlearns an output style that dramatically shorter than full rollout, teacher, and even its own baseline model. We show the length of the 10 percentile, 50 percentile (media), 90 percentile and the average length. *Right:* How student’s output aligns with the teacher’s distribution modes. “Top-1 Prob” shows the probability of the top 1 token’s probability mass, indicating how much the student commits in its first choice. “=tch. top-1” means the percentage of times where the student’s first choice is same as that of the teacher. “ $\in$ tch. top 2–5” means the percentage of times where the student’s first choice is the teachers 2nd to 5th choice. And “ $\notin$ tch. top-5” are the times the student’s first choice isn’t in the teacher’s top 5 list. Our method tends to choose teacher’s non-first choice and is more confident in its choice than the full rollout version.

Model	Response length (tokens)					Student commitment & mode alignment			
	10th ptile	50th ptile	90th ptile	Mean	Ratio	Top-1 prob	= tch. top-1	$\in$ tch. top 2–5	$\notin$ tch. top-5
Base student	~400	~860	~1,500	~990	2.2×	0.71	28.6%	59.8%	11.5%
Teacher (Q3-1.7B)	~700	~1,150	~1,800	~1,210	2.6×	—	—	—	—
Full Rollout	~1,190	~1,530	~1,770	~1,480	3.2×	0.77	45.7%	44.6%	9.7%
ESR	~190	~380	~800	~460	1.0×	0.79	41.9%	47.4%	10.7%

top-1 probability to 0.79 (vs 0.77 full-seq, 0.71 base) while *lowering* agreement with the teacher’s argmax to 41.9% (vs 45.7% full-seq) and raising the rate at which the student picks tokens in the teacher’s top-2 to top-5 to 47.4% (vs 44.6%). ESR commits harder, but on different teacher-supported tokens; full-seq commits toward the teacher’s argmax and inherits the teacher’s mistakes on those positions.

This is the mode-seeking mechanism in numbers. When one of those teacher-supported alternatives happens to be the correct mode, ESR exceeds the teacher; full-seq cannot, because supervising the late tail pulls the student toward the teacher’s argmax and the teacher’s mistakes along with it. A full proof of this hypothesis is out of scope; we note only that it is the natural reverse-KL-mode-seeking story applied to the planning region, and that it is consistent with every surpass-teacher cell we observe.

6 Related Work

Knowledge Distillation for Language Models. Knowledge distillation [Hinton et al., 2015, Kim and Rush, 2016] for autoregressive LLMs hinges on the choice of divergence and data distribution. Gu et al. [2024] argued for reverse KL, and Agarwal et al. [2024] introduced on-policy GKD, which supervises student rollouts with teacher logits. Subsequent work has explored skew KL with adaptive off-policy schedules [Ko et al., 2024], generalized f -divergences [Wen et al., 2023], mode- vs. mean-seeking dynamics [Wu et al., 2025], and speculative-decoding-style filtering [Xu et al., 2025].

We build on the on-policy reverse-KL paradigm but ask an orthogonal question: holding divergence and data distribution fixed, *which token positions along the student’s rollout carry useful signal?*

Token-Level Importance in Distillation and Reasoning. Not all tokens contribute equally to learning. Wang et al. [2025] showed that $\sim 20\%$ high-entropy “forking tokens” suffice for RL gradients; Vassoyan et al. [2025] relaxed KL on critical tokens to enable exploration; Li et al. [2025] identified a preplan-and-anchor mechanism in attention; and Singh and Hakkani-Tür [2026] pruned reasoning chains by functional importance. In distillation specifically, recent work selects tokens via teacher verification [Huang et al., 2025], adaptive temperature [Xie et al., 2026], student entropy along position/class/sample axes [Tavor et al., 2026], and entropy plus preference ranking [Kim and Baek, 2026]. Our Section 5.1 ablation contradicts the shared premise: selecting tokens by top-KL, various top-entropy variants, soft-weighted joint entropy, or top- K -renorm RKL¹ all underperform ESR and most underperform plain full-sequence training. Position is the load-bearing axis; scalar saliencies are not. This also reconciles Wang et al. [2025] with our findings: in on-policy distillation, high-entropy tokens are disproportionately concentrated in the uncontaminated early window.

Training Stability and Sequence Length. Full-sequence on-policy distillation is unstable: likelihood training degenerates into bland, repetitive text [Holtzman et al., 2020], and self-training drives model collapse [Shumailov et al., 2024]. Jang et al. [2026] proposed a geometric logit-space bridge to suppress harmful gradients on low-confidence tokens. On the length axis, Chen et al. [2025] showed that supervised training on the first 50% of tokens retains $\sim 91\%$ of full-sequence performance, and Yan et al. [2026] addressed exposure bias in long-CoT distillation via temperature scheduling. Our approach is complementary: restricting the loss scope to early positions, where teacher–student agreement is naturally highest, yields stability through a simpler mechanism and eliminates the degenerate repetition unique to on-policy training.

Curriculum Learning. Curriculum learning [Bengio et al., 2009] and scheduled sampling [Bengio et al., 2015] gradually increase task difficulty. For LLM pretraining, sequence-length warmup stabilizes training [Li et al., 2022] and variable-length curricula accelerate convergence [Pouransari et al., 2024]. Our progressive position schedule applies the curriculum principle to the distillation loss scope rather than to the input length—a form of curriculum specific to knowledge distillation.

Reasoning Distillation. Chain-of-thought prompting [Wei et al., 2022] made intermediate reasoning the unit of transfer, and DeepSeek-R1 [DeepSeek-AI, 2025] showed that SFT-distilling a strong reasoner into a smaller model is often competitive with direct RL. On reasoning tasks (MATH-500, HumanEval/MBPP, BFCL) and a 1.5B-math-student $\rightarrow$ 8B-teacher sweep, ESR surpasses the teacher on several cells, indicating that gains beyond the teacher are possible when the mode-seeking reverse-KL loss is restricted to the planning region—explaining why the on-policy, early-stopped-rollout regime is more efficient and stable than full-sequence on-policy distillation, which frequently degrades on coding and math.

7 Conclusion

In on-policy knowledge distillation, the teacher’s per-token supervision quality is a strong, monotonic function of position: early tokens carry strong, independent signal about reasoning strategy; late tokens, scored on the student’s own contaminated previous output, contribute weak, dependent signal that is actively harmful when trained on. A one-line intervention — restrict the loss and the rollout to the first N response tokens — yields four simultaneous wins: (i) **better**, matching or exceeding full-sequence distillation on every task/family/regime we test and surpassing the teacher in nine cells; (ii) **faster**, $20\text{--}35\times$ wall-clock with $4\times$ less GPU memory; (iii) **stabler**, eliminating HumanEval collapse, BFCL parse-rate collapse, and the MATH-500 boxed-repetition failure mode; and (iv) **structural**, in that selecting the same number of tokens by any KL or entropy saliency fails — position is an independent axis, not a proxy for scalar surprise. Teaching the student how to begin is sufficient to improve the entire response.

Limitations. Our experiments are conducted on small open-source models with a limited data budget; whether the ESR story holds at industrial scale (much larger students, teachers, and training

¹TODO: resolve and cite the top- K -renorm RKL method (Fu et al., 2025).

306 corpora) remains unclear. We also do not test multi-modality or long-horizon tasks, both of which
307 may exhibit different positional signal-quality patterns.

308 **Broader Impacts.** ESR reduces the wall-clock and memory cost of on-policy distillation by 20–
309 $35\times$ and $\sim 4\times$ respectively, which lowers the barrier to legitimate distillation research for academic
310 and small-industry groups that cannot afford full-rollout pipelines. The same efficiency gain has a
311 dual-use aspect: it lowers the cost of distilling capabilities from undisclosed or proprietary teacher
312 models, with potential implications for model-extraction and IP. We use only publicly released
313 open-weight teachers and standard public benchmarks in this work, but practitioners deploying ESR
314 on closed teachers should respect the teachers’ terms of use.

315 References

- 316 Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist,
317 and Olivier Bachem. On-policy distillation of language models: Learning from self-generated
318 mistakes. In *International Conference on Learning Representations*, 2024.
- 319 Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan,
320 Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language
321 models. *arXiv preprint arXiv:2108.07732*, 2021.
- 322 Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. Scheduled sampling for sequence
323 prediction with recurrent neural networks. In *Advances in Neural Information Processing Systems*,
324 2015.
- 325 Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In
326 *Proceedings of the 26th Annual International Conference on Machine Learning*, pages 41–48,
327 2009.
- 328 Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared
329 Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large
330 language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- 331 Wei-Rui Chen, Vignesh Kothapalli, Ata Fatahibaarzi, Hejian Sang, Shao Tang, Qingquan Song,
332 Zhipeng Wang, and Muhammad Abdul-Mageed. Distilling the essence: Efficient reasoning
333 distillation via sequence truncation. *arXiv preprint arXiv:2512.21002*, 2025.
- 334 Alex Cloud, Minh Le, James Chua, Jan Betley, Anna Sztyber-Betley, Jacob Hilton, Samuel Marks,
335 and Owain Evans. Subliminal learning: Language models transmit behavioral traits via hidden
336 signals in data. *arXiv preprint arXiv:2507.14805*, 2025.
- 337 DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning.
338 *arXiv preprint arXiv:2501.12948*, 2025.
- 339 Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large
340 language models. In *International Conference on Learning Representations*, 2024.
- 341 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
342 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In
343 *Advances in Neural Information Processing Systems*, 2021.
- 344 Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. In
345 *NIPS Deep Learning and Representation Learning Workshop*, 2015.
- 346 Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text
347 degeneration. In *International Conference on Learning Representations*, 2020.
- 348 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
349 and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International
350 Conference on Learning Representations*, 2022.
- 351 Haiduo Huang, Jiangcheng Song, Yadong Zhang, and Pengju Ren. SelecTKD: Selective token-
352 weighted knowledge distillation for LLMs. *arXiv preprint arXiv:2510.24021*, 2025.

353 Ijun Jang, Jewon Yeom, Juan Yeo, Hyunggu Lim, and Taesup Kim. Stable on-policy distillation
354 through adaptive target reformulation. *arXiv preprint arXiv:2601.07155*, 2026.

355 Minsang Kim and Seung Jun Baek. Explain in your own words: Improving reasoning via token-
356 selective dual knowledge distillation. *arXiv preprint arXiv:2603.13260*, 2026.

357 Yoon Kim and Alexander M. Rush. Sequence-level knowledge distillation. In *Proceedings of the*
358 *2016 Conference on Empirical Methods in Natural Language Processing*, pages 1317–1327, 2016.

359 Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. DistiLLM: Towards streamlined
360 distillation for large language models. In *International Conference on Machine Learning*, 2024.

361 Conglong Li, Minjia Zhang, and Yuxiong He. The stability-efficiency dilemma: Investigating
362 sequence length warmup for training GPT models. In *Advances in Neural Information Processing*
363 *Systems*, 2022.

364 Yang Li, Zhichen Dong, Yuhan Sun, Weixun Wang, Shaopan Xiong, Yijia Luo, Jiashun Liu, Han
365 Lu, Jiamang Wang, Wenbo Su, Bo Zheng, and Junchi Yan. Attention illuminates LLM reason-
366 ing: The preplan-and-anchor rhythm enables fine-grained policy optimization. *arXiv preprint*
367 *arXiv:2510.13554*, 2025.

368 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
369 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint*
370 *arXiv:2305.20050*, 2023.

371 Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by
372 ChatGPT really correct? rigorous evaluation of large language models for code generation. In
373 *Advances in Neural Information Processing Systems*, 2023.

374 Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Ji, Vishnu Suresh, Ion Stoica, and Joseph E.
375 Gonzalez. The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation
376 of large language models. In *Proceedings of the 42nd International Conference on Machine*
377 *Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48371–48392, 2025.

378 Hadi Pouransari, Chun-Liang Li, Jen-Hao Rick Chang, Pavan Kumar Anasosalu Vasu, Cem Koc,
379 Vaishaal Shankar, and Oncel Tuzel. Dataset decomposition: Faster LLM training with variable
380 sequence length curriculum. In *Advances in Neural Information Processing Systems*, 2024.

381 Qwen Team. Qwen2.5-math technical report: Toward mathematical expert model via self-
382 improvement. *arXiv preprint arXiv:2409.12122*, 2024.

383 Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

384 Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. AI
385 models collapse when trained on recursively generated data. *Nature*, 631(8022):755–759, 2024.

386 Janvijay Singh and Dilek Hakkani-Tür. Do LLMs encode functional importance of reasoning tokens?
387 *arXiv preprint arXiv:2601.03066*, 2026.

388 Almog Tavor, Itay Ebenspanger, Neil Cnaan, and Mor Geva. Rethinking selective knowledge
389 distillation. *arXiv preprint arXiv:2602.01395*, 2026.

390 Jean Vassoyan, Nathanael Beau, and Roman Plaud. Ignore the KL penalty! boosting exploration
391 on critical tokens to enhance RL fine-tuning. In *Findings of the Association for Computational*
392 *Linguistics: NAACL 2025*, 2025.

393 Shenzi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
394 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
395 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 rule: High-entropy minority tokens drive
396 effective reinforcement learning for LLM reasoning. In *Advances in Neural Information Processing*
397 *Systems*, 2025.

398 Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc
399 Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In
400 *Advances in Neural Information Processing Systems*, 2022.

- 401 Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f-divergence minimization for sequence-level
402 knowledge distillation. In *Proceedings of the 61st Annual Meeting of the Association for Computa-*
403 *tional Linguistics (Volume 1: Long Papers)*, pages 10817–10834, 2023.
- 404 Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. Rethinking
405 Kullback-Leibler divergence in knowledge distillation for large language models. In *Proceedings*
406 *of the International Conference on Computational Linguistics (COLING)*, 2025.
- 407 Xurong Xie, Zhucun Xue, Jiafu Wu, Jian Li, Yabiao Wang, Xiaobin Hu, Yong Liu, and Jiangn-
408 ing Zhang. LLM-oriented token-adaptive knowledge distillation. In *Proceedings of the AAAI*
409 *Conference on Artificial Intelligence*, 2026.
- 410 Wenda Xu, Rujun Han, Zifeng Wang, Long T. Le, Dhruv Madeka, Lei Li, William Yang Wang,
411 Rishabh Agarwal, Chen-Yu Lee, and Tomas Pfister. Speculative knowledge distillation: Bridging
412 the teacher-student gap through interleaved sampling. In *International Conference on Learning*
413 *Representations*, 2025.
- 414 Shaotian Yan, Kaiyuan Liu, Chen Shen, Bing Wang, Sinan Fan, Jun Zhang, Yue Wu, Zheng Wang,
415 and Jieping Ye. Distribution-aligned sequence distillation for superior long-CoT reasoning. *arXiv*
416 *preprint arXiv:2601.09088*, 2026.

417 A Training Efficiency

Table 5: **Training efficiency, detailed breakdown.** Per-step wall-clock time on a single A6000 (48GB). Student: Qwen2.5-Math-1.5B, LoRA, bs=16.

Teacher	Method	Gen (s)	Score (s)	Train (s)	Total (s)
Q3-1.7B	Full-seq	100–731	3–10	5–12	~280
	Ours ($N=100$)	~5	~1	~2	~8
Q3-4B	Full-seq	97–733	3–9	6–10	~210
	Ours ($N=100$)	~5	~1	~2	~8
Q3-8B	Full-seq	96–230	7–11	1–6	~170 [†]
	Ours ($N=100$)	~5	~1	~2	~8

[†] Frequent OOMs; 48GB insufficient for 8B teacher + vLLM + student.

418 B Additional Student-Teacher Pairs

Table 6: **Additional student-teacher pairs.** All use ESR-100, LoRA, 3,200 problems.

Student	Teacher	MATH-500 avg@4	BFCL full_acc	HumanEval pass@1
Qwen3-1.7B	— (baseline)	69.20%	59.80%	39.63%
Qwen3-1.7B	Qwen3-4B	69.20%	73.70%	44.51%
Qwen3-1.7B	Qwen3-8B	69.15%	74.00%	40.24%
Qwen3-4B	— (baseline)	77.95%	73.00%	73.17%
Qwen3-4B	Qwen3-8B	77.05%	76.80%	71.34%

419 C Full Experimental Results

420 C.1 Math Results: Per-Step Performance

421 Table 7 presents per-step results for the primary math experiments (LoRA, $n=1$, 3,200 problems).

Table 7: **MATH-500 per-step results.** LoRA, $n=1$, $bs=16$, 3,200 problems. Baseline: 50.95% avg@4.

Method	Metric	Step 50	Step 100	Step 150	Step 200
ESR-50	avg@4	62.35	66.05	66.65	64.85
	maj@4	69.40	72.00	71.00	71.20
	pass@4	77.20	79.40	81.00	79.60
ESR-100	avg@4	63.75	64.45	65.15	65.85
	maj@4	70.00	68.40	69.60	70.80
	pass@4	79.80	78.40	80.20	79.80
ESR-150	avg@4	65.35	66.65	65.30	65.75
	maj@4	66.80	67.00	66.30	67.30
	pass@4	79.00	81.00	78.20	80.00
ESR-200	avg@4	66.05	64.65	65.10	65.55
	maj@4	71.20	68.40	70.00	71.20
	pass@4	81.00	79.80	80.60	80.60

Table 8: **Complete MATH-500 results, LoRA, $n=16$, $bs=16$.** Best per configuration in **bold**.

Config	Step 50	Step 100	Step 150	Step 200
<i>avg@4</i>				
ESR-5	55.15	55.65	56.50	56.30
ESR-10	59.50	56.40	56.30	56.00
ESR-20	60.20	57.35	56.40	56.55
ESR-50	62.45	62.15	60.50	61.90
ESR-100	61.15	63.55	64.25	63.85
ESR-150	63.55	65.70	64.80	64.40
ESR-200	63.25	66.75	64.55	65.25
Progressive	61.10	62.05	61.95	62.30
Full-seq [†]	65.60	46.30	49.70	43.80
Full-seq (first-boxed)	65.55	64.95	65.55	64.95
<i>pass@4</i>				
ESR-5	74.0	74.8	76.0	73.4
ESR-10	75.8	73.2	75.4	74.2
ESR-20	77.0	76.0	73.4	75.8
ESR-50	77.0	78.4	77.6	78.6
ESR-100	77.6	79.2	80.2	80.2
ESR-150	79.6	77.4	80.0	79.4
ESR-200	79.6	80.2	79.4	79.2
Progressive	77.2	77.8	78.2	78.8
Full-seq	80.0	72.2	77.0	74.2

[†] Last-boxed extraction; degrades due to repetition.

Table 9: **Complete coding results, LoRA.** HumanEval (HE) and MBPP pass@1.

Config	Metric	s50	s100	s150	s200	s250	s300	s350	s400
ESR-50	HE	37.8	39.0	39.6	41.5	40.2	40.9	42.1	40.9
	MBPP	52.1	48.9	46.0	46.0	47.6	46.3	46.6	46.3
ESR-100	HE	37.2	39.0	42.1	37.8	39.0	37.8	37.8	38.4
	MBPP	52.1	49.7	49.2	49.2	48.4	49.2	47.9	48.9
ESR-150	HE	36.6	35.4	36.6	39.0	41.5	39.6	38.4	37.2
	MBPP	51.3	50.0	48.4	50.5	49.5	50.3	49.7	48.7
Full-seq	HE	40.2	31.7	32.3	32.9	27.4	28.0	26.8	26.8
	MBPP	52.6	48.9	49.5	48.1	47.9	47.1	48.4	47.6

422 C.2 Math Results: Earlier Series ($n=16$, 200 problems)

423 C.3 Coding Results

424 C.4 Function Calling Results

Table 10: **Function calling results (BFCL), LoRA.** Name accuracy / Full accuracy / Parse rate. Best full_acc in **bold**.

Method	Best Step	Name Acc	Full Acc	Parse Rate
Baseline	—	9.70%	2.70%	24.20%
Teacher (Qwen3-1.7B)	—	75.30%	54.00%	75.30%
ESR-50	200	95.20%	57.20%	98.30%
ESR-100	100	86.20%	61.30%	91.30%
ESR-150	200	88.70%	61.50%	92.50%
ESR-200	200	80.80%	54.50%	90.20%
Full-seq	100	81.00%	58.20%	86.70%

425 C.5 Scaling Experiment Results

Table 11: **Scaling results.** All use ESR-100, LoRA, 3,200 problems. Best across training steps. Baseline performance shown for reference.

Student → Teacher	Math avg@4	BFCL full_acc	HE pass@1	MBPP+ pass@1
M-1.5B baseline	50.95%	2.70%	—	—
M-1.5B → Q3-1.7B	65.85%	61.30%	42.07%	44.44%
M-1.5B → Q3-4B	68.95%	45.80%	43.29%	47.62%
M-1.5B → Q3-8B	67.85%	57.00%	41.46%	44.44%
Q3-1.7B baseline	69.20%	59.80%	39.63%	44.71%
Q3-1.7B → Q3-4B	69.20%	73.70%	44.51%	47.62%
Q3-1.7B → Q3-8B	69.15%	74.00%	40.24%	48.15%
Q3-4B baseline	77.95%	73.00%	73.17%	55.56%
Q3-4B → Q3-8B	77.05%	76.80%	71.34%	56.35%

426 C.6 Full-Sequence Degradation Details

427 Table 12 summarises every full-rollout cell (across pairs, tasks, and regimes) for which we have a
 428 per-step trajectory on disk and observe degradation across training. For each cell we report the *peak*
 429 score (best checkpoint across training) and the *lowest* score (worst checkpoint after the peak), together
 430 with the step indices at which they occur. Bold Δ values indicate cells where Full Rollout drops
 431 below the student baseline. The Qwen-math $n=16$ row (boxed-repetition collapse) is documented in
 432 greater detail in Table 13 below.

433 **Per-step trajectories** for the seven cells in Table 12:

- 434 • **Qwen-1.5B→Q3-1.7B, LoRA, HumanEval (pass@1):** 40.2 (s50) → 31.7 → 32.3 → 32.9
 435 → 27.4 → 28.0 → 26.8 → 26.8 (s400).
- 436 • **Qwen-1.5B→Q3-1.7B, LoRA, HumanEval+ (pass@1):** 35.4 (s50) → 27.4 → 29.9 →
 437 29.9 → 25.0 → 26.2 → 25.0 → 25.0 (s400).
- 438 • **Qwen-1.5B→Q3-1.7B, FullFT, HumanEval (pass@1):** 32.3 → 32.3 → 36.6 (s150) →
 439 36.0 → 31.1 → 30.5 → 31.7 → 31.1 (s400).
- 440 • **Qwen-1.5B→Q3-1.7B, FullFT, HumanEval+ (pass@1):** 27.4 → 28.7 → 31.7 (s150) →
 441 30.5 → 26.2 → 26.2 → 26.8 → 26.2 (s400).
- 442 • **Gemma-2-2B→Gemma-3-4B, LoRA, HumanEval (pass@1):** 22.0 (s50) → 14.6 → 11.6
 443 → 11.6 (s200). Falls below the 23.78 student baseline at s100 and stays there.

Table 12: **Full-rollout collapses across training (peak $\rightarrow$ lowest).** Each row is a single Full Rollout experiment; we report the highest and lowest checkpoint scores across all saved steps. “ Δ ” is lowest $-$ peak. Cells with † end the run *below the student baseline*. Trajectories are taken from the per-step eval logs on the training cluster (docs/main_results.md, docs/scaling_results.md). ESR on the same cells (not shown here) stays within ± 1 pp of its peak across training; see Tables 7 and 9.

Pair	Regime	Task	Metric	Peak (step)	Lowest (step)	Δ
Qwen-1.5B $\rightarrow$ Q3-1.7B	LoRA	MATH-500 ($n=16$)	avg@4	65.60 (s50)	43.80 (s200) †	−21.80
Qwen-1.5B $\rightarrow$ Q3-4B	LoRA	MATH-500	avg@4	67.45 (s50)	54.45 (s100)	−13.00
Qwen-1.5B $\rightarrow$ Q3-4B	LoRA	MATH-500	maj@4	73.20 (s50)	63.40 (s200)	−9.80
Qwen-1.5B $\rightarrow$ Q3-1.7B	LoRA	HumanEval	pass@1	40.2 (s50)	26.8 (s350)	−13.4
Qwen-1.5B $\rightarrow$ Q3-1.7B	LoRA	HumanEval+	pass@1	35.4 (s50)	25.0 (s250)	−10.4
Qwen-1.5B $\rightarrow$ Q3-1.7B	FullIFT	HumanEval	pass@1	36.6 (s150)	30.5 (s300)	−6.1
Qwen-1.5B $\rightarrow$ Q3-1.7B	FullIFT	HumanEval+	pass@1	31.7 (s150)	26.2 (s250)	−5.5
Gemma-2-2B $\rightarrow$ Gemma-3-4B	LoRA	HumanEval	pass@1	22.0 (s50)	11.6 (s150) †	−10.4

† Lowest score is below the student baseline (Qwen-Math-1.5B MATH-500: 50.95%; Gemma-2-2B HumanEval: 23.78%).

- **Qwen-1.5B $\rightarrow$ Q3-4B, LoRA, MATH-500 avg@4:** 67.45 (s50) $\rightarrow$ 54.45 $\rightarrow$ 58.85 $\rightarrow$ 55.05 (s200) (boxed-repetition pattern; ESR-100 on the same pair climbs to 68.95).
- **Qwen-1.5B $\rightarrow$ Q3-1.7B, LoRA, MATH-500 ($n=16$) avg@4:** 65.60 (s50) $\rightarrow$ 46.30 $\rightarrow$ 49.70 $\rightarrow$ 43.80 (s200). Boxed-repetition collapse; statistics in Table 13.

Table 13: **Boxed repetition progression** in full-sequence LoRA distillation on MATH-500 ($n=16$, bs=16). Companion to the corresponding row of Table 12.

Statistic	Step 50	Step 100	Step 150	Step 200
Multi-boxed responses	2.9%	91.5%	80.5%	89.1%
Avg <code>\boxed{}</code> per response	1.0	88.1	54.1	58.0
Repetitive responses	5.9%	78.8%	58.3%	77.5%
Avg response length (chars)	1,388	4,023	3,585	4,558
Official avg@4	65.60%	46.30%	49.70%	43.80%
Relaxed avg@4 (GT in text)	74.60%	70.90%	71.90%	70.30%

D KL-Profile Diagnostic for Choosing N

A practical question for a new student–teacher pair is: what value of N should I use? Our position-sweep (Section 4.3) shows a wide plateau on Qwen math ($N \in [50, 200]$, avg@4 within ~ 1 pp); the same plateau appears on Gemma math ($N \approx 100$). These plateau locations coincide with the position at which the cumulative per-position KL $\sum_{k \leq N} \mathbb{E}[\text{KL}_k]$ reaches $\sim 45\%$ of the full-trajectory total, measured on a small (~ 1 k trajectory) pre-training rollout. This suggests a cheap rule of thumb: run a 1k-trajectory per-position KL scan before training, and pick N at the 45% cumulative-KL point. We verify this on two cells (Qwen, Gemma math) and leave a full six-cell verification ($\{\text{Qwen, Gemma}\} \times \{\text{math, coding, funcall}\}$) as a short follow-up (**TBD: KL-profile scan**); the diagnostic is reported here as a practical recipe rather than a claim.

E Token Classification Methodology

We classify each token into six categories based on string matching:

1. **planning:** Reasoning keywords (“To”, “Let”, “First”, “Step”, “We”, “Given”, “Therefore”, “Thus”, “Since”).
2. **structural:** Punctuation, whitespace, formatting tokens.
3. **math_number:** Digits (0–9).
4. **math_operator:** Arithmetic operators (+, −, $\times$, $/$, =).

465

5. **math_latex**: LaTeX delimiters ($\backslash \langle$, $\backslash \rangle$).

466

6. **continuation**: All others.

Table 14: Mean KL by token category and position range.

Category	0–4	5–19	20–49	50–99	100–199	200–499
planning	4.50	0.79	1.49	1.66	1.49	2.37
structural	3.26	1.46	1.60	0.93	0.60	0.86
math_number	1.49	0.60	0.74	0.28	0.17	0.13
math_operator	7.30	1.84	0.81	0.37	0.21	0.14
math_latex	8.84	9.23	6.50	4.95	2.97	1.87
continuation	1.94	1.12	1.19	0.89	0.70	0.48

Table 15: Top 20 highest-KL tokens (minimum 50 occurrences across 10,000 trajectories).

Rank	Token	Category	Count	Mean KL
1	“Solution”	planning	152	21.93
2	“Analysis”	continuation	125	16.49
3	$\backslash \langle$	math_latex	7,152	13.21
4	“examines”	continuation	74	11.51
5	“He”	continuation	150	10.80
6	$\backslash \langle$	math_latex	21,243	10.30
7	“First”	planning	1,706	9.98
8	“tests”	continuation	52	8.71
9	$\backslash \backslash$	math_latex	82	8.68
10	“There”	continuation	201	8.28
11	“Therefore”	planning	4,913	7.95
12	“Identify”	continuation	1,345	7.78
13	“To”	planning	8,806	6.90
14	“When”	continuation	89	6.62
15	“This”	continuation	621	6.34
16	“Thus”	planning	2,547	6.24
17	“First” (space)	planning	259	6.10
18	“To” (space)	planning	1,174	5.37
19	“The”	planning	2,626	5.25
20	“Next”	planning	1,753	5.03

467 F Assets and Licenses

468 Table 16 lists the models and datasets used in this paper, with their providers and licenses. All assets
 469 are used in accordance with their respective terms of use.

Table 16: Assets used in this paper. All licenses verified at time of submission.

Asset	Type	Provider	License
Qwen2.5-Math-1.5B	Model	Alibaba	Apache 2.0
Qwen2.5-Math-7B	Model	Alibaba	Apache 2.0
Qwen3-1.7B / 4B / 8B / 14B	Model	Alibaba	Apache 2.0
Gemma-2-2B	Model	Google	Gemma Terms of Use
Gemma-3-4B	Model	Google	Gemma Terms of Use
NuminaMath	Dataset	Numina	Apache 2.0
CodeUltraFeedback	Dataset	Coseal	MIT
glaive-function-calling-v2	Dataset	Glaive AI	Apache 2.0
MATH-500	Benchmark	Hendrycks et al.	MIT
HumanEval / HumanEval+	Benchmark	OpenAI / EvalPlus	MIT
BFCL	Benchmark	UC Berkeley	Apache 2.0

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and §1 state that ESR matches or exceeds full-rollout distillation while being $20\text{--}35\times$ faster with $\sim 4\times$ less memory; each claim maps to Table 1, §4.3, §4.4, and §5.1.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: A *Limitations* paragraph at the end of §7 discusses the small open-source-model regime, limited data budget, and untested industrial scale.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[N/A\]](#)

Justification: The paper is empirical; no formal theorems or proofs.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: §4 lists student/teacher pairs, datasets, all hyperparameters, the loss, and evaluation protocols; Appendices A and C give per-step timing and results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [\[No\]](#)

Justification: We will release the training code and LoRA adapters publicly after acceptance; the methodology in §4 is detailed enough to reproduce the main results in the meantime.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: All training and evaluation settings are stated in the Setup paragraph of §4, with per-step results in Appendix C.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [\[No\]](#)

Justification: Reported numbers are single-seed best-step results; robustness comes from breadth across families, tasks, and teacher scales rather than per-cell variance, and this is acknowledged in §7.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

519 Answer: [Yes]
520 Justification: Appendix A (Table 5) reports per-step wall-clock on a single A6000 (48GB)
521 for each teacher scale, broken down by phase.

522 **9. Code of ethics**

523 Question: Does the research conducted in the paper conform, in every respect, with the
524 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

525 Answer: [Yes]
526 Justification: We use open-weight models and public benchmarks, with no human subjects,
527 private data, or scraped content.

528 **10. Broader impacts**

529 Question: Does the paper discuss both potential positive societal impacts and negative
530 societal impacts of the work performed?

531 Answer: [Yes]
532 Justification: A *Broader Impacts* paragraph at the end of §7 discusses both the positive impact
533 of cheaper distillation research and the dual-use risk of lower-cost capability extraction from
534 proprietary teachers.

535 **11. Safeguards**

536 Question: Does the paper describe safeguards that have been put in place for responsible
537 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
538 image generators, or scraped datasets)?

539 Answer: [N/A]
540 Justification: No new pretrained base models or scraped datasets are released; trained
541 students are LoRA adapters on top of public bases.

542 **12. Licenses for existing assets**

543 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
544 the paper, properly credited and are the license and terms of use explicitly mentioned and
545 properly respected?

546 Answer: [Yes]
547 Justification: All models and datasets are cited in §4 and a consolidated license table is
548 provided in Appendix F.

549 **13. New assets**

550 Question: Are new assets introduced in the paper well documented and is the documentation
551 provided alongside the assets?

552 Answer: [N/A]
553 Justification: No new assets are released at submission; LoRA adapters and code will be
554 documented and released after acceptance.

555 **14. Crowdsourcing and research with human subjects**

556 Question: For crowdsourcing experiments and research with human subjects, does the paper
557 include the full text of instructions given to participants and screenshots, if applicable, as
558 well as details about compensation (if any)?

559 Answer: [N/A]
560 Justification: No crowdsourcing or human subjects.

561 **15. Institutional review board (IRB) approvals or equivalent for research with human**
562 **subjects**

563 Question: Does the paper describe potential risks incurred by study participants, whether
564 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
565 approvals (or an equivalent approval/review based on the requirements of your country or
566 institution) were obtained?

567 Answer: [N/A]

568 Justification: No human subjects research.

569 **16. Declaration of LLM usage**

570 Question: Does the paper describe the usage of LLMs if it is an important, original, or
571 non-standard component of the core methods in this research? Note that if the LLM is used
572 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
573 scientific rigor, or originality of the research, declaration is not required.

574 Answer: [N/A]

575 Justification: LLMs are the object of study in this work and do not require declaration. We
576 additionally used LLMs as coding and writing assistants, but all methodological choices,
577 experimental design, analyses, and conclusions were author-led; the LLMs did not contribute
578 original methodology.