
Less is More: Early Stopping Rollout for On-Policy Distillation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 On-policy distillation scores every token of the student’s rollout with the teacher.
2 We identify a subtlety that changes the picture: the teacher’s score at position t is
3 conditioned on the student’s own previously generated tokens, so early in a rollout
4 the teacher judges *how to approach the problem*, whereas later it can only judge *how*
5 *to continue what the student has already written*. Given a sufficiently long student
6 rollout, the teacher’s accuracy collapses toward the student baseline, indicating
7 that late-position supervision is far weaker than early-position supervision. We
8 therefore propose *Early Stopping Rollout* (ESR): restrict both the rollout and the
9 reverse-KL loss to the first N response tokens. Across tasks, model families, LoRA
10 and full fine-tuning, and a range of teacher scales, ESR matches or exceeds full-
11 sequence distillation nearly everywhere, sometimes surpasses the teacher, trains
12 far faster with far less memory, and avoids the collapse modes of full-sequence
13 training. The obvious explanation—ESR as a proxy for high-KL or high-entropy
14 tokens—fails: every scalar-saliency selector we tried underperforms ESR, and
15 most underperform plain full-sequence. Position is an independent, load-bearing
16 axis—it is where the planning tokens live, and a short reverse-KL loss there is
17 enough to improve the entire response.

18 1 Introduction

19 On-policy distillation (OPD) has emerged as a dominant paradigm for model distillation in industrial
20 practice. The student generates its own rollouts τ , which are then scored by the teacher: at each token,
21 the teacher’s probability $\pi_{teacher}(\tau_{student} \mid x_{prompt})$ serves as the soft target for the student [Agarwal
22 et al., 2024, Gu et al., 2024]. Viewed through an RL lens, this is equivalent to using the teacher as a
23 dense, token-level reward model that judges the student’s own behavior on a given prompt. The final
24 objective is then a simple *uniform average* of the per-token reverse-KL across the trajectory.

25 This uniform averaging implicitly assumes that the teacher’s signal is equally reliable at every position
26 in the trajectory. We argue that this assumption is unfounded, and that it breaks down precisely
27 because of the autoregressive structure of LLMs. At the first token, the teacher’s score is conditioned
28 only on the prompt: $\pi_{teacher}(\tau_{student}^{t=1} \mid x_{prompt})$ - indeed what we expect the teacher to score on. At
29 a later position m , however, it becomes conditioned also on the student’s own previously generated
30 tokens: $\pi_{teacher}(\tau_{student}^{t=m} \mid x_{prompt}, \tau_{student}^{t=1:m})$. This conditioning context is off-policy to the teacher
31 model and may even be out-of-distribution (OOD), and grows more so with position.

32 To probe how severe this drift is, we run a preliminary experiment that measures how the teacher
33 behaves when continuing from an early-stopped student rollout. As shown in Figure 1, the teacher’s
34 quality collapses toward the student’s within roughly 300 tokens. If the teacher’s late-position signal
35 is unreliable, the natural question is whether early tokens alone are systematically more informative
36 — and whether training on them is sufficient.

Motivated by this, we propose Early Stopping Rollout(ESR): restrict the student rollout to its first N tokens (e.g., 50 or 100) and compute the distillation loss only on this early window. The change is a single line in any on-policy distillation loop. Despite its simplicity, ESRconsistently outperforms full-rollout OPD across tasks (math, code, function calling), training regimes (LoRA, full fine-tuning), model scales (students 1.5B–7B, teachers 1.7B–14B), and model families (Qwen, Gemma), while reducing wall-clock cost by up to $24\times$ and peak training memory by up to $4\times$. We further observe that full-rollout OPD frequently destabilizes and collapses during training, whereas ESRdoes not collapse in any of our settings. Perhaps most surprisingly, although the student is normally upper-bounded by the teacher’s performance, ESR-trained students often *exceed* the teacher.

Such a small change yields such large gains probe us to investigate deeper. We have verified that later position tokens suffer from signal quality issues. But we want to further validate whether the effectiveness of earlier tokens is intrinsically caused by better KL and entropy signals, because we also find that early tokens generally have larger KL and entropy. We test selectors that pick tokens by high KL, high entropy or combined regardless of position; all underperform ESR, with most underperforming full-rollout. Therefore, although we don’t exclude KL and entropy as auxillary factors, we confirm that the position is a an independent factor rather a saliency proxy. The late position tokens have signal quality issues in themselves. Furthermore, we identified two other mechanisms that may be working together: (i) *Strategy learning*: we find content-wise the early tokens tend to encode strategies and plans, while later tokens tend to focus on specific executions. It may be better to learn just the planning part from the teacher and leave the execution to its own way. (ii) *Convergence Cascade*: after training on the early window, KL divergence on the *untrained* late tokens also drops by 30–40%; the KL divergence does not have to see late positions to repair them.

These mechanisms explain why ESRcan beat full-rollout. They do not explain why ESRsometimes even *exceeds* the teacher — a stronger phenomenon we observe on most of our main experiment Table 1. We trace this to the mode-seeking nature of reverse-KL: the student commits to a subset of supported teacher modes instead of chasing all of them. Our ESRloss helps it steer toward a non-dominant but correct mode rather than the teacher’s argmax. Empirically, we find ESR-trained students commit harder to their own choices than full-rollout students yet agree *less* with the teacher’s argmax, and produce trajectories $\sim 3\times$ shorter than full-rollout and $\sim 2\times$ shorter than the base student (§5.3). It learns the teacher’s strategy but preserves its succinct style, leading to better compound behaviors. This finding showcases a potential path of superceding the teacher model in distillation that worth further investigations.

Our contributions are summarized below:


1. **Early Stopping Rollout(ESR) outperforms full-rollout on-policy distillation.** A one-line change—restricting the rollout length to the first N response tokens—beats full-sequence distillation across tasks, model families, scales, and training regimes, while being dramatically more efficient and far more stable to train. It can even exceed the teacher’s own performance (Sections 3–4.3).
2. **A systematic investigation of why it works.** First, through a token-selection ablation, we show that the gain is not “high KL” or “high entropy” in disguise: position itself is a load-bearing axis. We also identify strategy learning, and convergence cascading to be potential factors. Furthermore, we explain why ESRoften surpasses the teacher: by exploiting the mode-seeking behavior of reverse-KL training, ESRpushes the student toward a better, non-dominant sub-mode (Section 5).

2 Why Late-Position Supervision Is Unreliable

In on-policy distillation, the teacher model scores the student’s rollout at each position t via $\pi_t(\cdot | x, y_{<t}^{\text{student}})$ and averages these scores uniformly across the rollout. The conditioning context $y_{<t}^{\text{student}}$, however, is whatever this particular student happened to write at this particular moment of training. As t grows it drifts away from any distribution the teacher itself was trained on—off-policy to the teacher and, for long enough or atypical enough rollouts, plausibly out-of-distribution.

To verify the intuition, we feed the teacher a k -token student-generated rollout on MATH-500 and letting it continue autoregressively, the teacher’s avg@4 ($n=4$, $t=0.7$) falls from its 65.30% unconditional baseline to 62.70% at $k=100$ and 51.75% at $k=300$ —essentially student-baseline

Figure 1: **The teacher loses accuracy quickly as the student-generated rollout grows in a few hundred of tokens.** MATH-500, avg@4 ($n=4$, $t=0.7$). Teacher = Qwen3-1.7B; student = Qwen2.5-Math-1.5B. After ~ 300 student tokens the teacher has effectively been dragged down to student-baseline performance.



(Figure 1). Late-position teacher scores are not independent assessments of the prompt; it is a guess at how to continue a trajectory the student has already committed to.

The problem compounds during training because the student’s policy is moving. Early on, the student’s rollout is far from the teacher; late-position targets are fitted under a strongly off-policy prefix. As training progresses and earlier tokens approach the teacher’s distribution, the right late-position behavior changes too—continuations that were sensible above an off-policy prefix are no longer sensible above an on-policy one. But the late-position habits learned earlier do not cleanly erase: gradient descent gives no guarantee of unlearning, and in LoRA’s low-rank setting it updates make redirection compete for capacity that may already be spent. Full-rollout training thus chases a moving late-position target while carrying ballast from earlier ones. This insight expectes training stability issues for full-rollout OPD – a prediction we later verify in §4.3

3 Our Method: Early Stopping Rollout (ESR) On Policy Distillation

3.1 Formulation

Let π_s denote the student and π_t the teacher. In standard on-policy reverse-KL distillation, the student generates a response $\mathbf{y} = (y_1, \dots, y_T)$ conditioned on prompt x , and the loss is

$$\mathcal{L}_{\text{full}} = \mathbb{E}_{\mathbf{y} \sim \pi_s(\cdot|x)} \left[\sum_{t=1}^T \text{KL}(\pi_s(\cdot | x, \mathbf{y}_{<t}) \parallel \pi_t(\cdot | x, \mathbf{y}_{<t})) \right]. \quad (1)$$

ESR (position cutoff N , with $N \ll T$ in practice) truncates the student rollout to its first N tokens, and the loss is computed over exactly those tokens:

$$\mathcal{L}_{\text{ESR}}(N) = \mathbb{E}_{\mathbf{y} \sim \pi_s(\cdot|x), |\mathbf{y}| \leq N} \left[\sum_{t=1}^{\min(N, |\mathbf{y}|)} \text{KL}(\pi_s(\cdot | x, \mathbf{y}_{<t}) \parallel \pi_t(\cdot | x, \mathbf{y}_{<t})) \right]. \quad (2)$$

If the student emits EOS before position N , the rollout terminates naturally. Everything else—generation temperature, LoRA target modules, optimizer, scorer—is unchanged from the standard on-policy KD loop.

Table 1: **Main results: ESR dominates Full Rollout across tasks and families.** Best score across training steps; metrics defined in §4 setup. 🟡: ESR surpasses the teacher reference. **Full Rollout** numbers mark the best step before degradation; ↓ flags cells where Full Rollout further decays with training (more details in Appendix C).

Pair	Method	MATH-500		Coding		BFCL
		avg@4	pass@4	HE	HE+	full_acc
Qwen2.5-Math 1.5B → Qwen3 1.7B	Student baseline	50.95	72.80	31.10	26.20	2.70
	Teacher reference	65.30	77.00	39.60	36.00	54.00
	Full Rollout	62.35↓	74.60↓	40.20↓	35.40↓	58.20🟡
	ESR	65.85🟡	79.80🟡	42.10🟡	36.00	61.30🟡
Gemma-2 2B → Gemma-3 4B	Student baseline	13.45	28.20	23.78	20.10	73.17
	Teacher reference	66.60	74.80	20.70	20.10	72.83
	Full Rollout	22.95	31.40	22.00↓	17.70↓	76.83🟡
	ESR	27.20	39.40	28.70🟡	24.40🟡	79.00🟡

4 Analysis on the Advancement of ESR

4.1 Setup

Models. We evaluate across multiple student-teacher pairs. Our primary student is Qwen2.5-Math-1.5B [Qwen Team, 2024], paired with teachers of increasing scale: Qwen3-1.7B, Qwen3-4B, and Qwen3-8B [Qwen Team, 2025]. We also tested Qwen2.5-Math-7B with Qwen3-14B to test student model size generality. To test cross-family generality, we also evaluate with gemma-2-2b-it (student) and gemma-3-4b-it (teacher).

Training. All experiments use LoRA [Hu et al., 2022] ($r=32$, $\alpha=64$, all linear layers, $\text{lr} = 5 \times 10^{-5}$) with reverse KL divergence loss. Temperature 0.7 for on-policy generation. Each training step processes a batch of 16 problems with 1 rollout per problem ($n_{\text{samples}}=1$, batch size 16). We train for 200 steps on all tasks, and saving checkpoints every 50 steps. Training data are drawn from NuminaMath, CodeUltraFeedback, and glaive-function-calling-v2. Our method uses $N=100$ unless otherwise specified.

Evaluation. MATH-500 [Hendrycks et al., 2021, Lightman et al., 2023] with $n=4$ samples at temperature 0.7 (reporting avg@4), HumanEval/HumanEval+ [Chen et al., 2021, Austin et al., 2021, Liu et al., 2023] at temperature 0.0 (reporting pass@1), BFCL [Patil et al., 2025] with 600 problems reporting full accuracy: correct function name and arguments.

4.2 Overall performance

Table 1 presents the main results across tasks and model families. Table 2 presents ablation results with full finetuning. Table 3 shows model size scaling results. We report the best performance achieved across all training steps for each configuration. But notice that full rollout distillation frequently collapse to much lower performance as it trains. We denote the sessions that collapses with a down arrows.

ESRbeats full rollout across model families and tasks, and frequently surpasses teacher performance. Across the cells of Table 1, ESR beats full rollout on every cell, and in six cells ESR surpasses the teacher outright (Qwen math avg@4, pass@4; Qwen HumanEval and HumanEval+; Qwen BFCL full_acc; Gemma coding on all four columns; Gemma BFCL full_acc). The Gemma coding results have the biggest gap: full rollout drives HumanEval pass@1 to 22.00, below the 23.78 student baseline and worse than the 20.70 teacher, while ESR recovers 31.10. Qwen-math full rollout is marked ↓ because its best number (62.35 avg@4) is a non-stable best step: under $n=16$ it collapses to 46.3% by step 100 (Appendix C.6).

ESRmatches or beats full rollout across full fine-tuning and LoRA settings. Table 2 reports FullFT on MATH-500 for both Qwen and Gemma. On Qwen **Full Rollout** 58.20 is slightly better than ours ESR 56.20 avg@4, but the gap is close. For Gemma ESR still dominates **Full Rollout**


Table 2: **FullFT on MATH-500**. Best across training steps. **Bold**: ESR beats **Full Rollout**.

Pair	Method	avg@4	pass@4
Qwen2.5-Math 1.5B	student/teacher	50.95/65.30	72.80/77.00
→ Qwen3 1.7B	Full Rollout	58.20	75.40
	ESR	56.20	73.80
Gemma-2 2B	student/teacher	13.45/66.60	28.20/74.80
→ Gemma-3 4B	Full Rollout	13.90	25.00
	ESR	26.65	40.40

Table 3: **Scaling teacher and student (LoRA) on MATH-500**. $\downarrow$: **Full Rollout** peak before degradation. **Bold**: ESR beats **Full Rollout**.

Pair	Method	avg@4	pass@4
Qwen2.5-Math 1.5B	student/teacher	50.95/77.95	72.80/86.40
→ Qwen3 4B	Full Rollout	67.45 $\downarrow$	80.60
	ESR	68.95	81.00
Qwen2.5-Math 1.5B	student/teacher	50.95/77.70	72.80/87.80
→ Qwen3 8B	Full Rollout	66.75	81.60
	ESR	67.85	82.20
Qwen2.5-Math 7B	student/teacher	53.60/76.15	75.00/83.20
→ Qwen3 14B	Full Rollout	68.85 $\downarrow$	80.00
	ESR	68.95	81.20

Figure 2: **Rollout Length N Sweep on MATH-500**. LoRA, Qwen2.5-Math-1.5B → Qwen3-1.7B; best avg@4 across training steps. Full rollout and the undistilled baseline are shown as horizontal references. Performance saturates for $N \in [50, 200]$ and all beat full-seq.



144 by +12.75pp avg@4 and +15.40pp pass@4. ESR is therefore the safer choice in both parameter
145 regimes.

146 **ESRscales across model sizes.** The advantage persists as both teacher and student grow. Table 3
147 reports the results scaling both the teacher and the student. With the same student (Qwen2.5-Math-
148 1.5B), we tried Qwen3-4B and 8B as the teacher. With Qwen3-4B, **Full Rollout** peaks at (67.45) and
149 then *collapses* to 55.05 (Appendix C.6), while ESR stays stable and climbs to 68.95. Against the 8B
150 teacher, ESR still edges **Full Rollout** (67.85 vs 66.75). The pattern persists at larger scale along with
151 student model too: with a Qwen2.5-Math-7B student and Qwen3-14B teacher, **Full Rollout** again
152 peaks at step 50 (68.85) and drifts to 62.40, while ESR holds stable across all checkpoints.

153 4.3 Stability of ESR

154 **ESRis more robust than full rollout in training.** Across our (pair $\times$ task $\times$ regime) matrix,
155 full-rollout distillation degrades with training very frequently; ESR degrades nowhere. Appendix
156 Table 14 reports each collapsing cell as a (peak $\rightarrow$ lowest) pair with the corresponding step indices.
157 For example, Qwen-LoRA coding falls from 40.2 to 26.8; Qwen-LoRA math collapses from 65.60 at
158 s50 to 43.80 ($\Delta = -21.80$ pp, ending below the 50.95% student baseline); Gemma-LoRA collapse
159 from 22.0 to 11.6 ($\Delta = -10.4$ pp). Three of seven full-rollout collapses end below the student
160 baseline. ESR generally hold a fluctuation within ± 3 pp of its peak across training (Tables 9, 11).

161 **Early Stopping Rollouts is not sensitive to the choice of N .** A natural question naturally occurs -
162 how to choose where to stop? Is it sensitive? We conducted a set of sweeping experiments in Figure 2
163 sweeps N on MATH-500 and reveals a robust region: it reaches just as good performance starting
164 from $N=50$ and remains stable to $N=200$. The method is not sensitive to the exact choice of N
165 within a certain region.

Table 4: **Training efficiency.** Single A6000 (48 GB), bs=16, student teacher Qwen3-1.7B. We report the average running time and GPU memory usage across student model generation, training and teacher scoring phases. Note that the real time usage can be larger if there isn’t enough GPU to hold the student and teacher models together and requires model loading and unloading.

Phase	Per-step wall time			Peak GPU memory		
	ESR (100 tok.)	Full rollout	Speedup	ESR (100 tok.)	Full rollout	Savings
Teacher scoring	1 s	7 s	7×	7.3 GB	14.9 GB	2.0×
Student generation	5 s	180 s	36×	7.2 GB	8.9 GB	1.2×
Student training	2 s	7 s	3.5×	9.6 GB	39.5 GB	4.1×
Total (single step)	8 s	194 s	24×	24.1 GB	63.3 GB	2.6×
Total (200 steps)	27 min	10.8 hr	24×			

4.4 Efficiency of ESR

Table 4 shows that ESR achieves a 24× wall-clock speedup and reduces peak training memory by ~ 4×. The dominant cost in full-seq is autoregressive generation (180 s/step for sequences averaging ~1000 tokens); ESR generates only $N=100$ tokens (5 s/step). Note that with ESR, all the student and teacher models can be put in one A6000 GPU comfortably. In our own practice, it saves a further big time overhead of model loading and unloading that we do not report here.

5 Why Does ESR Work?

The result that a one-line change matches or exceeds full-sequence distillation — at a fraction of the compute, with stability everywhere the full-sequence baseline degrades — is surprising enough that it demands a deeper mechanism. We structure this section around three questions: (§5.1) Is the early-position effect simply a proxy for high-KL or high-entropy tokens? (§5.2) If not, what additional structural reasons make the early window privileged? (§5.3) Why does ESR sometimes *exceed* the teacher?

5.1 Is the Effectiveness of the Early Position Caused by KL and Entropy?

Figure 3 shows that early positions are simultaneously high-KL and high-entropy, so a natural alternative to positional selection is *information-theoretic* selection: pick the k tokens with the highest KL divergence or highest student/teacher entropy, regardless of position. We tested every such selector (Table 5): all underperform ESR, and most underperform full-sequence. For each strategy in Table 5 we keep the full 400-token on-policy rollout, compute the selector, and apply the reverse-KL loss only to the selected tokens.

Two findings fall out of Table 5. (i) Top-RKL=100, the selector that captures the most raw KL signal, is the *worst* (53.35%) — worse than no distillation. Increasing k to 200 only partly recovers. (ii) Every KL/entropy/combined selector underperforms full-sequence (62.35%), which in turn underperforms ESR (65.85%). These selectors all concentrate on the contaminated late region where the teacher is scoring its own continuation of the student’s previously generated tokens; the benefit of the early region is therefore *structural*, not reducible to any scalar per-token statistic.

Takeaway. ESR is not mainly caused by KL or entropy. Position is an independent, load-bearing axis — not “high-KL in disguise” and not “high-entropy in disguise”. Whatever makes the early window privileged is structural; we identify two such reasons next.

5.2 Strategy Learning and Convergence Cascade

Section 5.1 rules out scalar per-token saliency. Beyond the late-position contamination already established in §2, we identify two further structural mechanisms that jointly explain why training on the early window is so effective.

Figure 3: **Early positions carry all three — student entropy, teacher entropy, and reverse-KL — simultaneously** (Qwen, MATH-500). Natural hypothesis: ESR works because those regions carry the most signal (high KL) at the most uncertain points (high entropy). If that were the mechanism, selecting tokens by high-KL or high-entropy alone should recover the gain. Table 5 shows it does not.




Table 5: **Token selection strategies on MATH-500.** Qwen2.5-Math-1.5B $\rightarrow$ Qwen3-1.7B, LoRA, $n_{\text{samples}}=1$, best avg@4 across training steps. All selectors operate on full-length rollouts; ESR is the only one that additionally truncates the rollout. Baseline (no distillation) 50.95%; full-seq 62.35%. H_s : student entropy; H_t : teacher entropy; RKL: reverse KL.

Selection method	k	Best avg@4
ESR-100 (ours)	100	65.85
Full-seq (all tokens)	—	62.35
Top-RKL	100	53.35
Top-RKL	200	56.15
Top- H_t (teacher entropy)	100	63.30
Top- H_s (student entropy)	100	62.70
Top- H_s	200	53.20
RKL $\cdot H_s$	100	56.90
$H_t \cdot H_s$ (product)	100	55.35
RKL $\cdot H_t \cdot H_s$ (triple product)	100	57.90

Figure 4: **Case study: early tokens encode strategy; later tokens merely execute.** A representative MATH-500 trajectory. The first 100 tokens establish the geometric setup, name the unknown, and identify the key relationship (altitude bisects the leg)—these are the choices that determine whether the rollout will succeed. The last 100 tokens carry out an algebraic computation that any solver can finish once the strategy is fixed. ESR supervises the first window; the second follows for free.

Problem. In an isosceles right triangle, the altitude to the hypotenuse has length $4\sqrt{2}$. What is the area?


First 100 tokens (strategy). “Let’s start by understanding the problem. In an isosceles right triangle, the two legs are equal, and the hypotenuse is $\sqrt{2}$ times the leg. If we denote each leg by a , then hypotenuse $= a\sqrt{2}$. The altitude to the hypotenuse is half the leg, because the altitude bisects the...”

Last 100 tokens (execution). “... altitude from the right-angle vertex to the hypotenuse. This altitude has length $h = \frac{a \cdot a}{a\sqrt{2}} = \frac{a^2}{a\sqrt{2}} = \frac{a}{\sqrt{2}} = \frac{a\sqrt{2}}{2}$. Setting this equal to $4\sqrt{2}$: $\frac{a\sqrt{2}}{2} = 4\sqrt{2}$, so $a\sqrt{2} = 8\sqrt{2}$, giving $a = 8$. The area is $\frac{1}{2} \times 8 \times 8 = \boxed{32}$.”

199 **Strategy learning — the early window is where planning tokens live.** The early-position KL
200 spike concentrates on planning / structural tokens: the highest-KL tokens in our corpus are “To” (KL
201 6.90, $n=8806$), “First” (9.98), “Let”, “The”, “Therefore” (7.95), “Thus”, along with format-opening
202 tokens $\backslash$, $\backslash$, and code-fence ‘‘‘ (see Table 17). These tokens *select* the reasoning trajectory —
203 different openings commit the student to algebra-with-variables, factoring, or numeric-estimation
204 paths on the same problem. The teacher’s conditional distribution at these tokens is genuinely
205 multi-modal (entropy 1.07), so reverse-KL has room to move the student to a single teacher-supported
206 mode. After the plan is committed, the remainder is execution: the teacher is near-deterministic
207 (entropy 0.14) and there is little signal to extract. *Teach the plan; the execution follows* (Figure 4)
208 — and no scalar saliency can recover the early-window region because what matters is where in the
209 rollout a token lives, not how surprising it is locally. Moreover, recently Cloud et al. [2025] shows
210 student models may be able to learn the teacher’s deep internal preference even with random numbers
211 generated by the teacher, called “subliminal learning”. Therefore, the early tokens may suffice to
212 inject the subliminal mindset to the student.

213 **Convergence cascade — training the early window lowers KL everywhere.** A natural worry
214 is that ESR helps only inside the trained window. The opposite holds. After ESR-200 training,

Figure 5: **The convergence cascade.** Per-position KL between distilled student and teacher, before vs. after ESR-200 training. Yellow band: positions $[0, 200]$ that actually receive training loss. Blue band: the KL gap closed by training. Positions 200+ — which see no direct training signal — drop to the same KL as the trained region, confirming that alignment on the early window cascades through the autoregressive rollout.



per-position KL between distilled student and teacher drops by 50–60% inside $[0, N]$ and by 30–40% outside it (Figure 5). Comparing distilled vs. base responses on the same prompts, late tokens — which receive zero direct supervision — change as much as early tokens in surface content (Jaccard to baseline = 0.251 for both halves). Autoregressive generation explains the cascade: once the student has committed to the right plan in the first N tokens, the remaining rollout is a lower-KL, more on-policy completion for the teacher, so no gradient is needed there. The loss does not have to see late positions to align them.

5.3 Why Does ESR Sometimes Exceed the Teacher?

ESR exceeds the teacher on several cells (Qwen math; Qwen and Gemma function calling; Gemma coding, where the teacher is in fact weaker than the student). We offer a hypothesis. Reverse KL $\text{KL}(\pi_s \parallel \pi_t)$ is mode-seeking: it penalizes the student for putting mass on tokens the teacher does not support, but not for concentrating mass on a single supported token. When the teacher is multi-modal at the planning positions—one mode being a verbose, wasteful plan and another a concise correct plan—a short early-window loss can force the student to commit to the smaller, correct mode rather than averaging (Figure 6). On function-calling this effect is pronounced: ESR-100 lifts Qwen BFCL full_acc from 54.0% (teacher) to 61.3%, and the chained Qwen3-4B→8B configuration raises it to 76.8% against a 73.0% teacher. Full-sequence training cannot achieve this because its late-position loss pulls the student toward the average teacher behaviour across all supported modes, undoing the concentration.

Behavioral signature. The hypothesis predicts that ESR-trained students should (a) produce shorter, less-hedging rollouts than the teacher, and (b) commit hard, but on *non-argmax* teacher modes. Table 6 confirms both. ESR-100 rollouts are shorter than the base student, the teacher, and full-seq across the entire length distribution — even the top-decile (~ 800 tokens) is below the base student’s median (~ 860). On the top-10% highest-KL positions ($n=110,894$), ESR raises student top-1 probability to 0.79 (vs 0.77 full-seq, 0.71 base) while *lowering* agreement with the teacher’s argmax to 41.9% (vs 45.7% full-seq) and raising the rate at which the student picks tokens in the teacher’s top-2 to top-5 to 47.4% (vs 44.6%). ESR commits harder, but on different teacher-supported tokens; full-seq commits toward the teacher’s argmax and inherits the teacher’s mistakes on those positions.

This is the mode-seeking mechanism in numbers. When one of those teacher-supported alternatives happens to be the correct mode, ESR exceeds the teacher; full-seq cannot, because supervising the

Figure 6: **Mode-seeking on the planning region.** Schematic of reverse-KL behaviour at a multi-modal planning position. The teacher supports two modes (e.g. a verbose plan and a concise correct plan); reverse KL penalises student mass outside the support but not concentration within it. A short early-window loss can collapse the student onto the better supported mode, allowing the student to exceed the teacher’s average behaviour. Full-sequence training reverts the student toward the averaged teacher across late positions, undoing the concentration.




Table 6: **Behavioral signature of ESR.** *Left:* Response length distribution on MATH-500 ($n=2,000$ rollouts per model; tokens estimated as chars/3.5). ESR-100 is shorter than base, teacher, and full-seq across the whole distribution: even its long tail (90th pctile ~ 800 tok) is shorter than the base student’s typical response (50th pctile ~ 860 tok). **Ratio** = mean length relative to ESR. Full-seq mean is below its median because length is clipped at max_new_tokens. *Right:* Top-10% of token positions ranked by $KL[\pi_s \parallel \pi_t]$ on 100 base-student trajectories ($n=110,894$). “= tch. top-1”, “ $\in$ tch. top 2–5”, and “ $\notin$ tch. top-5” are mutually exclusive and sum to 100% (teacher omitted because the comparison is student-vs-teacher). ESR-100 commits hardest (top-1 prob 0.79) but agrees *less* with the teacher’s argmax than full-seq, and picks teacher-supported alternatives more often.

Model	Response length (tokens)					Student commitment & mode alignment			
	10th pctile	50th pctile	90th pctile	Mean	Ratio	Top-1 prob	= tch. top-1	$\in$ tch. top 2–5	$\notin$ tch. top-5
Base student	~ 400	~ 860	$\sim 1,500$	~ 990	$2.2\times$	0.71	28.6%	59.8%	11.5%
Teacher (Q3-1.7B)	~ 700	$\sim 1,150$	$\sim 1,800$	$\sim 1,210$	$2.6\times$	—	—	—	—
Full-seq	$\sim 1,190$	$\sim 1,530$	$\sim 1,770$	$\sim 1,480$	$3.2\times$	0.77	45.7%	44.6%	9.7%
ESR-100 (ours)	~ 190	~ 380	~ 800	~ 460	$1.0\times$	0.79	41.9%	47.4%	10.7%

late tail pulls the student toward the teacher’s argmax and the teacher’s mistakes along with it. A full proof of this hypothesis is out of scope; we note only that it is the natural reverse-KL-mode-seeking story applied to the planning region, and that it is consistent with every surpass-teacher cell we observe.

6 Related Work

Knowledge Distillation for Language Models. Knowledge distillation [Hinton et al., 2015] transfers knowledge from a teacher to a smaller student via soft targets. Kim and Rush [2016] extended this to sequence models with word-level and sequence-level objectives. For autoregressive LLMs, the choice of divergence and data distribution is decisive. Gu et al. [2024] argued for reverse KL in generative LLM distillation, on the grounds that reverse KL avoids the student placing probability on low-support teacher regions; Agarwal et al. [2024] introduced Generalized Knowledge Distillation (GKD), which generates rollouts from the student (“on-policy”) and supervises with teacher logits—substantially outperforming supervised off-policy distillation on reasoning tasks. Ko et al. [2024] proposed skew KL and an adaptive off-policy schedule; Wen et al. [2023] generalized sequence-level KD to arbitrary f -divergences; Wu et al. [2025] revisited the mode- vs. mean-seeking interpretation of reverse and forward KL in the LLM setting and showed that both converge under sufficient training but differ in which distribution regions they first concentrate on. Xu et al. [2025] introduced Speculative Knowledge Distillation, which interleaves teacher-student sampling in the style of speculative decoding and filters tokens the teacher would not have produced. Our work builds directly on the on-policy reverse-KL paradigm of Gu et al. [2024] and Agarwal et al. [2024], but

asks a different question: holding the divergence and data distribution fixed, *which token positions along the student’s rollout carry useful signal?* Our answer—only the early window (Early Stopping Rollout)—is orthogonal to the divergence-design axis these works explore.

Token-Level Importance in Distillation and Reasoning. A growing body of work demonstrates that not all tokens contribute equally to learning. Wang et al. [2025] showed that only $\sim 20\%$ of chain-of-thought tokens have high entropy, and these “forking tokens” steer reasoning paths—applying RL gradients to only these tokens matches full-gradient performance. In the RL fine-tuning setting, Vassoyan et al. [2025] demonstrated that uniform KL penalties across all token positions prevent exploration on “critical tokens” where the pretrained model is most uncertain, and proposed entropy-weighted KL relaxation that allows the model to deviate from its prior precisely where it needs to learn. Li et al. [2025] discovered a “preplan-and-anchor” mechanism in LLM attention patterns, where models first perform long-range contextual reference before producing semantic anchor tokens, enabling targeted credit assignment. Singh and Hakkani-Tür [2026] proposed greedy pruning of reasoning chains and showed that LLMs internally encode token-level functional importance for answer generation.

In the distillation setting specifically, several concurrent works have explored token selection. Huang et al. [2025] proposed SelecTKD, a propose-and-verify framework where the teacher verifies student-proposed tokens, applying full loss to accepted tokens and masking rejected ones. Xie et al. [2026] introduced AdaKD, which adaptively adjusts token-level temperature based on learning stability. Tavor et al. [2026] systematically disentangled selective KD along position, class, and sample axes, proposing student-entropy-guided position selection (SE-KD) that reduces wall time by 70% without sacrificing accuracy. Kim and Baek [2026] combined entropy-based token selection with preference ranking in a dual distillation framework.

These works collectively establish that not all tokens deserve equal gradient weight, and most of them pick the “important” tokens via a scalar per-token criterion—entropy, KL, difficulty, or a verification score. Our Section 5.1 ablation directly contradicts the premise that such scalars are sufficient: selecting the same number of tokens by top-KL, top-entropy (student, teacher, AND, OR, triple-product), soft-weighted joint entropy, or top- K -renorm RKL (a vocabulary-level renormalization of RKL proposed for stabilizing distillation; exact citation to be added in camera-ready)¹ all underperform ESR and most underperform plain full-sequence training. Position is the load-bearing axis; scalar token saliencies are not. This also resolves a tension with Wang et al. [2025], who find that RL gradients concentrated on high-entropy “forking tokens” suffice: in on-policy *distillation*, the high-entropy tokens are disproportionately in the uncontaminated early window, so their result is consistent with ours once this early-window concentration of entropy is accounted for.

Training Stability and Sequence-Level Issues. Full-sequence on-policy distillation is known to suffer from instability. Holtzman et al. [2020] identified that likelihood-based training can produce bland, repetitive text—a form of neural text degeneration. Shumailov et al. [2024] demonstrated that iteratively training generative models on their own outputs leads to model collapse, where distribution tails disappear. In on-policy KD specifically, Jang et al. [2026] proposed Veto, a geometric bridge in logit space that suppresses harmful gradients on low-confidence tokens, directly addressing the instability we also observe. Our approach is complementary: rather than modifying the target distribution, we restrict the loss scope to early positions where teacher-student agreement is naturally highest, achieving stability through a simpler mechanism.

Several works have studied how sequence length affects distillation quality. Chen et al. [2025] showed that training on only the first 50% of tokens in supervised reasoning distillation retains $\sim 91\%$ of full-sequence performance, establishing a computation-quality tradeoff. Yan et al. [2026] addressed exposure bias and capacity misalignment in long chain-of-thought distillation through temperature-scheduled learning. Our work extends the truncation insight to the on-policy setting: we demonstrate that positional loss masking on student-generated sequences achieves equal or better performance than full-sequence loss, while providing the additional benefit of eliminating degenerate repetition patterns that are unique to on-policy training.

Curriculum Learning and Sequence Scheduling. Curriculum learning [Bengio et al., 2009] improves training by gradually increasing task difficulty. Bengio et al. [2015] applied this principle

¹TODO: resolve and cite the top- K -renorm RKL method (Fu et al., 2025).

to sequence models through scheduled sampling, transitioning from teacher-forced to free-running generation during training. In the LLM pretraining setting, Li et al. [2022] showed that long sequences at training start cause extreme gradient variance, and sequence length warmup stabilizes training while enabling larger batch sizes. Pouransari et al. [2024] extended this with variable sequence length curriculum for LLM pretraining, achieving up to $3\times$ faster convergence. Our progressive position schedule—linearly increasing the position limit from 1 to 200 over training—applies the curriculum principle to the distillation loss scope rather than to the input sequence length, and represents a novel form of curriculum specific to knowledge distillation.

Reasoning Distillation. Chain-of-thought prompting [Wei et al., 2022] made intermediate reasoning the unit of transfer for distillation, and the DeepSeek-R1 release [DeepSeek-AI, 2025] established that SFT-distilling a strong reasoner into a smaller model is often competitive with direct RL on the small model. Our experiments include three reasoning tasks (MATH-500, HumanEval/MBPP, BFCL) and a 1.5B-math-student $\rightarrow$ 8B-teacher sweep; ESR surpasses the teacher on several of these cells (Qwen math, function-calling), indicating that reasoning-distillation gains beyond the teacher are possible when the mode-seeking reverse-KL loss is restricted to the planning region. This is a complementary perspective to the DeepSeek-R1 recipe, which uses fully-sequenced SFT: our analysis explains why the on-policy, early-stopped-rollout regime is both more efficient and more stable for reasoning-style chain-of-thought distillation, and why full-sequence on-policy distillation—which in principle should be the strongest setting—frequently degrades across training on coding and math.

7 Conclusion

In on-policy knowledge distillation, the teacher’s per-token supervision quality is a strong, monotonic function of position: early tokens carry strong, independent signal about reasoning strategy; late tokens, scored on the student’s own contaminated previous output, contribute weak, dependent signal that is actively harmful when trained on. A one-line intervention — restrict the loss and the rollout to the first N response tokens — yields four simultaneous wins: (i) **better**, matching or exceeding full-sequence distillation on every task/family/regime we test and surpassing the teacher in nine cells; (ii) **faster**, 20–35 $\times$ wall-clock with 4 $\times$ less GPU memory; (iii) **stabler**, eliminating HumanEval collapse, BFCL parse-rate collapse, and the MATH-500 boxed-repetition failure mode; and (iv) **structural**, in that selecting the same number of tokens by any KL or entropy saliency fails — position is an independent axis, not a proxy for scalar surprise. Teaching the student how to begin is sufficient to improve the entire response.

Limitations. Our experiments are conducted on small open-source models with a limited data budget; whether the ESR story holds at industrial scale (much larger students, teachers, and training corpora) remains unclear. We also do not test multi-modality or long-horizon tasks, both of which may exhibit different positional signal-quality patterns.

Broader Impacts. ESR reduces the wall-clock and memory cost of on-policy distillation by 20–35 $\times$ and $\sim 4\times$ respectively, which lowers the barrier to legitimate distillation research for academic and small-industry groups that cannot afford full-rollout pipelines. The same efficiency gain has a dual-use aspect: it lowers the cost of distilling capabilities from undisclosed or proprietary teacher models, with potential implications for model-extraction and IP. We use only publicly released open-weight teachers and standard public benchmarks in this work, but practitioners deploying ESR on closed teachers should respect the teachers’ terms of use.

References

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*, 2024.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.

368 Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. Scheduled sampling for sequence
369 prediction with recurrent neural networks. In *Advances in Neural Information Processing Systems*,
370 2015.

371 Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In
372 *Proceedings of the 26th Annual International Conference on Machine Learning*, pages 41–48,
373 2009.

374 Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared
375 Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large
376 language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.

377 Wei-Rui Chen, Vignesh Kothapalli, Ata Fatahibaarzi, Hejian Sang, Shao Tang, Qingquan Song,
378 Zhipeng Wang, and Muhammad Abdul-Mageed. Distilling the essence: Efficient reasoning
379 distillation via sequence truncation. *arXiv preprint arXiv:2512.21002*, 2025.

380 Alex Cloud, Minh Le, James Chua, Jan Betley, Anna Sztyber-Betley, Jacob Hilton, Samuel Marks,
381 and Owain Evans. Subliminal learning: Language models transmit behavioral traits via hidden
382 signals in data. *arXiv preprint arXiv:2507.14805*, 2025.

383 DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning.
384 *arXiv preprint arXiv:2501.12948*, 2025.

385 Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large
386 language models. In *International Conference on Learning Representations*, 2024.

387 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
388 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In
389 *Advances in Neural Information Processing Systems*, 2021.

390 Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. In
391 *NIPS Deep Learning and Representation Learning Workshop*, 2015.

392 Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text
393 degeneration. In *International Conference on Learning Representations*, 2020.

394 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
395 and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International
396 Conference on Learning Representations*, 2022.

397 Haiduo Huang, Jiangcheng Song, Yadong Zhang, and Pengju Ren. SelecTKD: Selective token-
398 weighted knowledge distillation for LLMs. *arXiv preprint arXiv:2510.24021*, 2025.

399 Ijun Jang, Jewon Yeom, Juan Yeo, Hyunggu Lim, and Taesup Kim. Stable on-policy distillation
400 through adaptive target reformulation. *arXiv preprint arXiv:2601.07155*, 2026.

401 Minsang Kim and Seung Jun Baek. Explain in your own words: Improving reasoning via token-
402 selective dual knowledge distillation. *arXiv preprint arXiv:2603.13260*, 2026.

403 Yoon Kim and Alexander M. Rush. Sequence-level knowledge distillation. In *Proceedings of the
404 2016 Conference on Empirical Methods in Natural Language Processing*, pages 1317–1327, 2016.

405 Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. DistiLLM: Towards streamlined
406 distillation for large language models. In *International Conference on Machine Learning*, 2024.

407 Conglong Li, Minjia Zhang, and Yuxiong He. The stability-efficiency dilemma: Investigating
408 sequence length warmup for training GPT models. In *Advances in Neural Information Processing
409 Systems*, 2022.

410 Yang Li, Zhichen Dong, Yuhao Sun, Weixun Wang, Shaopan Xiong, Yijia Luo, Jiashun Liu, Han
411 Lu, Jiamang Wang, Wenbo Su, Bo Zheng, and Junchi Yan. Attention illuminates LLM reason-
412 ing: The preplan-and-anchor rhythm enables fine-grained policy optimization. *arXiv preprint
413 arXiv:2510.13554*, 2025.

414 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
415 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint*
416 *arXiv:2305.20050*, 2023.

417 Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by
418 ChatGPT really correct? rigorous evaluation of large language models for code generation. In
419 *Advances in Neural Information Processing Systems*, 2023.

420 Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Ji, Vishnu Suresh, Ion Stoica, and Joseph E.
421 Gonzalez. The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation
422 of large language models. In *Proceedings of the 42nd International Conference on Machine*
423 *Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48371–48392, 2025.

424 Hadi Pouransari, Chun-Liang Li, Jen-Hao Rick Chang, Pavan Kumar Anasosalu Vasu, Cem Koc,
425 Vaishaal Shankar, and Orel Tuzel. Dataset decomposition: Faster LLM training with variable
426 sequence length curriculum. In *Advances in Neural Information Processing Systems*, 2024.

427 Qwen Team. Qwen2.5-math technical report: Toward mathematical expert model via self-
428 improvement. *arXiv preprint arXiv:2409.12122*, 2024.

429 Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

430 Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. AI
431 models collapse when trained on recursively generated data. *Nature*, 631(8022):755–759, 2024.

432 Janvijay Singh and Dilek Hakkani-Tür. Do LLMs encode functional importance of reasoning tokens?
433 *arXiv preprint arXiv:2601.03066*, 2026.

434 Almog Tavor, Itay Ebenspanger, Neil Cnaan, and Mor Geva. Rethinking selective knowledge
435 distillation. *arXiv preprint arXiv:2602.01395*, 2026.

436 Jean Vassoyan, Nathanael Beau, and Roman Plaud. Ignore the KL penalty! boosting exploration
437 on critical tokens to enhance RL fine-tuning. In *Findings of the Association for Computational*
438 *Linguistics: NAACL 2025*, 2025.

439 Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
440 Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen
441 Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 rule: High-entropy minority tokens drive
442 effective reinforcement learning for LLM reasoning. In *Advances in Neural Information Processing*
443 *Systems*, 2025.

444 Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc
445 Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In
446 *Advances in Neural Information Processing Systems*, 2022.

447 Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f-divergence minimization for sequence-level
448 knowledge distillation. In *Proceedings of the 61st Annual Meeting of the Association for Computa-*
449 *tional Linguistics (Volume 1: Long Papers)*, pages 10817–10834, 2023.

450 Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. Rethinking
451 Kullback-Leibler divergence in knowledge distillation for large language models. In *Proceedings*
452 *of the International Conference on Computational Linguistics (COLING)*, 2025.

453 Xurong Xie, Zhucun Xue, Jiafu Wu, Jian Li, Yabiao Wang, Xiaobin Hu, Yong Liu, and Jiangn-
454 ing Zhang. LLM-oriented token-adaptive knowledge distillation. In *Proceedings of the AAAI*
455 *Conference on Artificial Intelligence*, 2026.

456 Wenda Xu, Rujun Han, Zifeng Wang, Long T. Le, Dhruv Madeka, Lei Li, William Yang Wang,
457 Rishabh Agarwal, Chen-Yu Lee, and Tomas Pfister. Speculative knowledge distillation: Bridging
458 the teacher-student gap through interleaved sampling. In *International Conference on Learning*
459 *Representations*, 2025.

460 Shaotian Yan, Kaiyuan Liu, Chen Shen, Bing Wang, Sinan Fan, Jun Zhang, Yue Wu, Zheng Wang,
461 and Jieping Ye. Distribution-aligned sequence distillation for superior long-CoT reasoning. *arXiv*
462 *preprint arXiv:2601.09088*, 2026.

Table 7: **Training efficiency, detailed breakdown.** Per-step wall-clock time on a single A6000 (48GB). Student: Qwen2.5-Math-1.5B, LoRA, bs=16. [†] Frequent OOMs; 48GB insufficient for 8B teacher + vLLM + student.

Teacher	Method	Gen (s)	Score (s)	Train (s)	Total (s)
Q3-1.7B	Full-seq	100–731	3–10	5–12	~280
	Ours ($N=100$)	~5	~1	~2	~8
Q3-4B	Full-seq	97–733	3–9	6–10	~210
	Ours ($N=100$)	~5	~1	~2	~8
Q3-8B	Full-seq	96–230	7–11	1–6	~170 [†]
	Ours ($N=100$)	~5	~1	~2	~8

Table 8: **Additional student-teacher pairs.** All use ESR-100, LoRA, 3,200 problems.

Student	Teacher	MATH-500 avg@4	BFCL full_acc	HumanEval pass@1
Qwen3-1.7B	— (baseline)	69.20%	59.80%	39.63%
Qwen3-1.7B	Qwen3-4B	69.20%	73.70%	44.51%
Qwen3-1.7B	Qwen3-8B	69.15%	74.00%	40.24%
Qwen3-4B	— (baseline)	77.95%	73.00%	73.17%
Qwen3-4B	Qwen3-8B	77.05%	76.80%	71.34%

A Training Efficiency

B Additional Student-Teacher Pairs

C Full Experimental Results

C.1 Math Results: Per-Step Performance

Table 9 presents per-step results for the primary math experiments (LoRA, $n=1$, 3,200 problems).

C.2 Math Results: Earlier Series ($n=16$, 200 problems)

C.3 Coding Results

C.4 Function Calling Results

C.5 Scaling Experiment Results

C.6 Full-Sequence Degradation Details

Table 14 summarises every full-rollout cell (across pairs, tasks, and regimes) for which we have a per-step trajectory on disk and observe degradation across training. For each cell we report the *peak* score (best checkpoint across training) and the *lowest* score (worst checkpoint after the peak), together with the step indices at which they occur. Bold Δ values indicate cells where Full Rollout drops below the student baseline. The Qwen-math $n=16$ row (boxed-repetition collapse) is documented in greater detail in Table 15 below.

Per-step trajectories for the seven cells in Table 14:

- **Qwen-1.5B→Q3-1.7B, LoRA, HumanEval (pass@1):** 40.2 (s50) → 31.7 → 32.3 → 32.9 → 27.4 → 28.0 → 26.8 → 26.8 (s400).
- **Qwen-1.5B→Q3-1.7B, LoRA, HumanEval+ (pass@1):** 35.4 (s50) → 27.4 → 29.9 → 29.9 → 25.0 → 26.2 → 25.0 → 25.0 (s400).
- **Qwen-1.5B→Q3-1.7B, FullFT, HumanEval (pass@1):** 32.3 → 32.3 → 36.6 (s150) → 36.0 → 31.1 → 30.5 → 31.7 → 31.1 (s400).

Table 9: **MATH-500 per-step results.** LoRA, $n=1$, $bs=16$, 3,200 problems. Baseline: 50.95% avg@4.

Method	Metric	Step 50	Step 100	Step 150	Step 200
ESR-50	avg@4	62.35	66.05	66.65	64.85
	maj@4	69.40	72.00	71.00	71.20
	pass@4	77.20	79.40	81.00	79.60
ESR-100	avg@4	63.75	64.45	65.15	65.85
	maj@4	70.00	68.40	69.60	70.80
	pass@4	79.80	78.40	80.20	79.80
ESR-150	avg@4	65.35	66.65	65.30	65.75
	maj@4	66.80	67.00	66.30	67.30
	pass@4	79.00	81.00	78.20	80.00
ESR-200	avg@4	66.05	64.65	65.10	65.55
	maj@4	71.20	68.40	70.00	71.20
	pass@4	81.00	79.80	80.60	80.60

Table 10: **Complete MATH-500 results, LoRA**, $n=16$, $bs=16$. Best per configuration in **bold**. [†] Last-boxed extraction; degrades due to repetition.

Config	Step 50	Step 100	Step 150	Step 200
<i>avg@4</i>				
ESR-5	55.15	55.65	56.50	56.30
ESR-10	59.50	56.40	56.30	56.00
ESR-20	60.20	57.35	56.40	56.55
ESR-50	62.45	62.15	60.50	61.90
ESR-100	61.15	63.55	64.25	63.85
ESR-150	63.55	65.70	64.80	64.40
ESR-200	63.25	66.75	64.55	65.25
Progressive	61.10	62.05	61.95	62.30
Full-seq [†]	65.60	46.30	49.70	43.80
Full-seq (first-boxed)	65.55	64.95	65.55	64.95
<i>pass@4</i>				
ESR-5	74.0	74.8	76.0	73.4
ESR-10	75.8	73.2	75.4	74.2
ESR-20	77.0	76.0	73.4	75.8
ESR-50	77.0	78.4	77.6	78.6
ESR-100	77.6	79.2	80.2	80.2
ESR-150	79.6	77.4	80.0	79.4
ESR-200	79.6	80.2	79.4	79.2
Progressive	77.2	77.8	78.2	78.8
Full-seq	80.0	72.2	77.0	74.2

- **Qwen-1.5B→Q3-1.7B, FullFT, HumanEval+ (pass@1):** 27.4 → 28.7 → 31.7 (s150) → 30.5 → 26.2 → 26.2 → 26.8 → 26.2 (s400).
- **Gemma-2-2B→Gemma-3-4B, LoRA, HumanEval (pass@1):** 22.0 (s50) → 14.6 → 11.6 → 11.6 (s200). Falls below the 23.78 student baseline at s100 and stays there.
- **Qwen-1.5B→Q3-4B, LoRA, MATH-500 avg@4:** 67.45 (s50) → 54.45 → 58.85 → 55.05 (s200) (boxed-repetition pattern; ESR-100 on the same pair climbs to 68.95).
- **Qwen-1.5B→Q3-1.7B, LoRA, MATH-500 ($n=16$) avg@4:** 65.60 (s50) → 46.30 → 49.70 → 43.80 (s200). Boxed-repetition collapse; statistics in Table 15.

D KL-Profile Diagnostic for Choosing N

A practical question for a new student–teacher pair is: what value of N should I use? Our position-sweep (Section 4.3) shows a wide plateau on Qwen math ($N \in [50, 200]$, avg@4 within $\sim 1pp$);

Table 11: **Complete coding results, LoRA.** HumanEval (HE) and MBPP pass@1.

Config	Metric	s50	s100	s150	s200	s250	s300	s350	s400
ESR-50	HE	37.8	39.0	39.6	41.5	40.2	40.9	42.1	40.9
	MBPP	52.1	48.9	46.0	46.0	47.6	46.3	46.6	46.3
ESR-100	HE	37.2	39.0	42.1	37.8	39.0	37.8	37.8	38.4
	MBPP	52.1	49.7	49.2	49.2	48.4	49.2	47.9	48.9
ESR-150	HE	36.6	35.4	36.6	39.0	41.5	39.6	38.4	37.2
	MBPP	51.3	50.0	48.4	50.5	49.5	50.3	49.7	48.7
Full-seq	HE	40.2	31.7	32.3	32.9	27.4	28.0	26.8	26.8
	MBPP	52.6	48.9	49.5	48.1	47.9	47.1	48.4	47.6

Table 12: **Function calling results (BFCL), LoRA.** Name accuracy / Full accuracy / Parse rate. Best full_acc in **bold**.

Method	Best Step	Name Acc	Full Acc	Parse Rate
Baseline	—	9.70%	2.70%	24.20%
Teacher (Qwen3-1.7B)	—	75.30%	54.00%	75.30%
ESR-50	200	95.20%	57.20%	98.30%
ESR-100	100	86.20%	61.30%	91.30%
ESR-150	200	88.70%	61.50%	92.50%
ESR-200	200	80.80%	54.50%	90.20%
Full-seq	100	81.00%	58.20%	86.70%

the same plateau appears on Gemma math ($N \approx 100$). These plateau locations coincide with the position at which the cumulative per-position KL $\sum_{k \leq N} \mathbb{E}[\text{KL}_k]$ reaches $\sim 45\%$ of the full-trajectory total, measured on a small ($\sim 1\text{k}$ trajectory) pre-training rollout. This suggests a cheap rule of thumb: run a 1k-trajectory per-position KL scan before training, and pick N at the 45% cumulative-KL point. We verify this on two cells (Qwen, Gemma math) and leave a full six-cell verification ($\{\text{Qwen, Gemma}\} \times \{\text{math, coding, funcall}\}$) as a short follow-up (**TBD: KL-profile scan**); the diagnostic is reported here as a practical recipe rather than a claim.

E Token Classification Methodology

We classify each token into six categories based on string matching:

1. **planning**: Reasoning keywords (“To”, “Let”, “First”, “Step”, “We”, “Given”, “Therefore”, “Thus”, “Since”).
2. **structural**: Punctuation, whitespace, formatting tokens.
3. **math_number**: Digits (0–9).
4. **math_operator**: Arithmetic operators (+, −, ×, /, =).
5. **math_latex**: LaTeX delimiters ($\backslash$, $\backslash$).
6. **continuation**: All others.

F Assets and Licenses

Table 18 lists the models and datasets used in this paper, with their providers and licenses. All assets are used in accordance with their respective terms of use.

Table 13: **Scaling results.** All use ESR-100, LoRA, 3,200 problems. Best across training steps. Baseline performance shown for reference.

Student → Teacher	Math avg@4	BFCL full_acc	HE pass@1	MBPP+ pass@1
M-1.5B baseline	50.95%	2.70%	—	—
M-1.5B → Q3-1.7B	65.85%	61.30%	42.07%	44.44%
M-1.5B → Q3-4B	68.95%	45.80%	43.29%	47.62%
M-1.5B → Q3-8B	67.85%	57.00%	41.46%	44.44%
Q3-1.7B baseline	69.20%	59.80%	39.63%	44.71%
Q3-1.7B → Q3-4B	69.20%	73.70%	44.51%	47.62%
Q3-1.7B → Q3-8B	69.15%	74.00%	40.24%	48.15%
Q3-4B baseline	77.95%	73.00%	73.17%	55.56%
Q3-4B → Q3-8B	77.05%	76.80%	71.34%	56.35%

Table 14: **Full-rollout collapses across training (peak → lowest).** Each row is a single Full Rollout experiment; we report the highest and lowest checkpoint scores across all saved steps. “ Δ ” is lowest – peak. Cells with † end the run *below the student baseline* (Qwen-Math-1.5B MATH-500: 50.95%; Gemma-2-2B HumanEval: 23.78%). Trajectories are taken from the per-step eval logs on the training cluster (docs/main_results.md, docs/scaling_results.md). ESR on the same cells (not shown here) stays within ± 1 pp of its peak across training; see Tables 9 and 11.

Pair	Regime	Task	Metric	Peak (step)	Lowest (step)	Δ
Qwen-1.5B → Q3-1.7B	LoRA	MATH-500 ($n=16$)	avg@4	65.60 (s50)	43.80 (s200) †	−21.80
Qwen-1.5B → Q3-4B	LoRA	MATH-500	avg@4	67.45 (s50)	54.45 (s100)	−13.00
Qwen-1.5B → Q3-4B	LoRA	MATH-500	maj@4	73.20 (s50)	63.40 (s200)	−9.80
Qwen-1.5B → Q3-1.7B	LoRA	HumanEval	pass@1	40.2 (s50)	26.8 (s350)	−13.4
Qwen-1.5B → Q3-1.7B	LoRA	HumanEval+	pass@1	35.4 (s50)	25.0 (s250)	−10.4
Qwen-1.5B → Q3-1.7B	FullIFT	HumanEval	pass@1	36.6 (s150)	30.5 (s300)	−6.1
Qwen-1.5B → Q3-1.7B	FullIFT	HumanEval+	pass@1	31.7 (s150)	26.2 (s250)	−5.5
Gemma-2-2B → Gemma-3-4B	LoRA	HumanEval	pass@1	22.0 (s50)	11.6 (s150) †	−10.4

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and §1 state that ESR matches or exceeds full-rollout distillation while being 20–35 $\times$ faster with $\sim 4\times$ less memory; each claim maps to Table 1, §4.3, §4.4, and §5.1.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: A *Limitations* paragraph at the end of §7 discusses the small open-source-model regime, limited data budget, and untested industrial scale.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is empirical; no formal theorems or proofs.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Table 15: **Boxed repetition progression** in full-sequence LoRA distillation on MATH-500 ($n=16$, $bs=16$). Companion to the corresponding row of Table 14.

Statistic	Step 50	Step 100	Step 150	Step 200
Multi-boxed responses	2.9%	91.5%	80.5%	89.1%
Avg <code>\boxed{}</code> per response	1.0	88.1	54.1	58.0
Repetitive responses	5.9%	78.8%	58.3%	77.5%
Avg response length (chars)	1,388	4,023	3,585	4,558
Official avg@4	65.60%	46.30%	49.70%	43.80%
Relaxed avg@4 (GT in text)	74.60%	70.90%	71.90%	70.30%

Table 16: **Mean KL by token category and position range.**

Category	0–4	5–19	20–49	50–99	100–199	200–499
planning	4.50	0.79	1.49	1.66	1.49	2.37
structural	3.26	1.46	1.60	0.93	0.60	0.86
math_number	1.49	0.60	0.74	0.28	0.17	0.13
math_operator	7.30	1.84	0.81	0.37	0.21	0.14
math_latex	8.84	9.23	6.50	4.95	2.97	1.87
continuation	1.94	1.12	1.19	0.89	0.70	0.48

Answer: [Yes]

Justification: §4 lists student/teacher pairs, datasets, all hyperparameters, the loss, and evaluation protocols; Appendices A and C give per-step timing and results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: We will release the training code and LoRA adapters publicly after acceptance; the methodology in §4 is detailed enough to reproduce the main results in the meantime.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: All training and evaluation settings are stated in the Setup paragraph of §4, with per-step results in Appendix C.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Reported numbers are single-seed best-step results; robustness comes from breadth across families, tasks, and teacher scales rather than per-cell variance, and this is acknowledged in §7.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix A (Table 7) reports per-step wall-clock on a single A6000 (48GB) for each teacher scale, broken down by phase.

Table 17: **Top 20 highest-KL tokens** (minimum 50 occurrences across 10,000 trajectories).

Rank	Token	Category	Count	Mean KL
1	“Solution”	planning	152	21.93
2	“Analysis”	continuation	125	16.49
3	\[	math_latex	7,152	13.21
4	“examines”	continuation	74	11.51
5	“He”	continuation	150	10.80
6	\(	math_latex	21,243	10.30
7	“First”	planning	1,706	9.98
8	“tests”	continuation	52	8.71
9	\\	math_latex	82	8.68
10	“There”	continuation	201	8.28
11	“Therefore”	planning	4,913	7.95
12	“Identify”	continuation	1,345	7.78
13	“To”	planning	8,806	6.90
14	“When”	continuation	89	6.62
15	“This”	continuation	621	6.34
16	“Thus”	planning	2,547	6.24
17	“First” (space)	planning	259	6.10
18	“To” (space)	planning	1,174	5.37
19	“The”	planning	2,626	5.25
20	“Next”	planning	1,753	5.03

Table 18: **Assets used in this paper.** All licenses verified at time of submission.

Asset	Type	Provider	License
Qwen2.5-Math-1.5B	Model	Alibaba	Apache 2.0
Qwen2.5-Math-7B	Model	Alibaba	Apache 2.0
Qwen3-1.7B / 4B / 8B / 14B	Model	Alibaba	Apache 2.0
Gemma-2-2B	Model	Google	Gemma Terms of Use
Gemma-3-4B	Model	Google	Gemma Terms of Use
NuminaMath	Dataset	Numina	Apache 2.0
CodeUltraFeedback	Dataset	Coseal	MIT
glaive-function-calling-v2	Dataset	Glaive AI	Apache 2.0
MATH-500	Benchmark	Hendrycks et al.	MIT
HumanEval / HumanEval+	Benchmark	OpenAI / EvalPlus	MIT
BFCL	Benchmark	UC Berkeley	Apache 2.0

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: We use open-weight models and public benchmarks, with no human subjects, private data, or scraped content.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: A *Broader Impacts* paragraph at the end of §7 discusses both the positive impact of cheaper distillation research and the dual-use risk of lower-cost capability extraction from proprietary teachers.

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

585 Answer: [N/A]

586 Justification: No new pretrained base models or scraped datasets are released; trained

587 students are LoRA adapters on top of public bases.

588 **12. Licenses for existing assets**

589 Question: Are the creators or original owners of assets (e.g., code, data, models), used in

590 the paper, properly credited and are the license and terms of use explicitly mentioned and

591 properly respected?

592 Answer: [Yes]

593 Justification: All models and datasets are cited in §4 and a consolidated license table is

594 provided in Appendix F.

595 **13. New assets**

596 Question: Are new assets introduced in the paper well documented and is the documentation

597 provided alongside the assets?

598 Answer: [N/A]

599 Justification: No new assets are released at submission; LoRA adapters and code will be

600 documented and released after acceptance.

601 **14. Crowdsourcing and research with human subjects**

602 Question: For crowdsourcing experiments and research with human subjects, does the paper

603 include the full text of instructions given to participants and screenshots, if applicable, as

604 well as details about compensation (if any)?

605 Answer: [N/A]

606 Justification: No crowdsourcing or human subjects.

607 **15. Institutional review board (IRB) approvals or equivalent for research with human**

608 **subjects**

609 Question: Does the paper describe potential risks incurred by study participants, whether

610 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

611 approvals (or an equivalent approval/review based on the requirements of your country or

612 institution) were obtained?

613 Answer: [N/A]

614 Justification: No human subjects research.

615 **16. Declaration of LLM usage**

616 Question: Does the paper describe the usage of LLMs if it is an important, original, or

617 non-standard component of the core methods in this research? Note that if the LLM is used

618 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

619 scientific rigor, or originality of the research, declaration is not required.

620 Answer: [N/A]

621 Justification: LLMs are the object of study in this work and do not require declaration. We

622 additionally used LLMs as coding and writing assistants, but all methodological choices,

623 experimental design, analyses, and conclusions were author-led; the LLMs did not contribute

624 original methodology.